

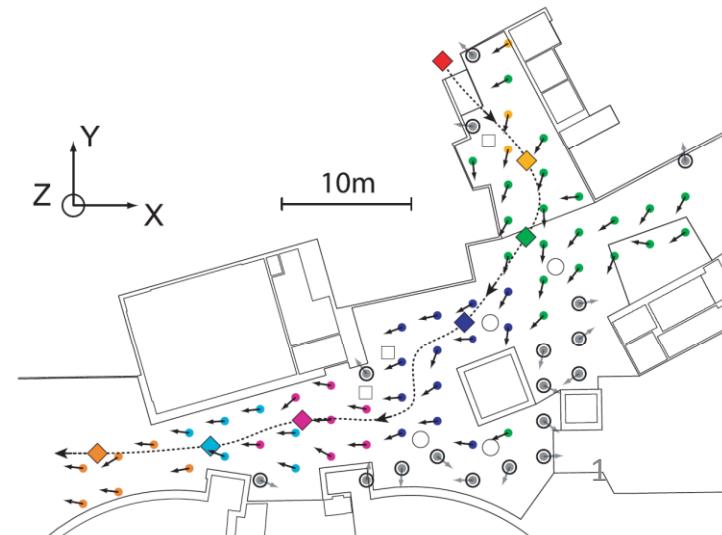
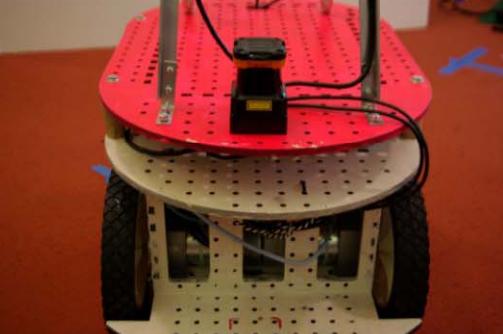
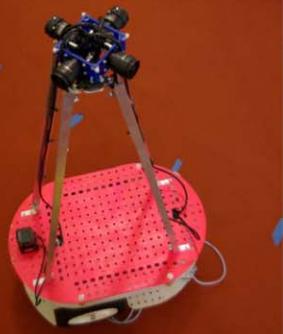
Body-Relative Navigation Guidance using Uncalibrated Cameras

Olivier Koch

PhD Thesis Defense

21-Jan-2010

Thesis Committee: Prof. Bill Freeman, Prof. Rob Miller, Prof. Seth Teller



Problem Statement

Human-oriented navigation in GPS-denied environments

- Finding your way in **unknown environments** with **no prior map** and **no external source of localization**



Soldiers in the field



Hazmat teams



Visually impaired

Problem Statement

Rely on internal sensing to provide navigation guidance:

- Accurate
- Spatially extended
- Temporally consistent
- Complex, cluttered, dynamic environments



Soldiers in the field



Hazmat teams

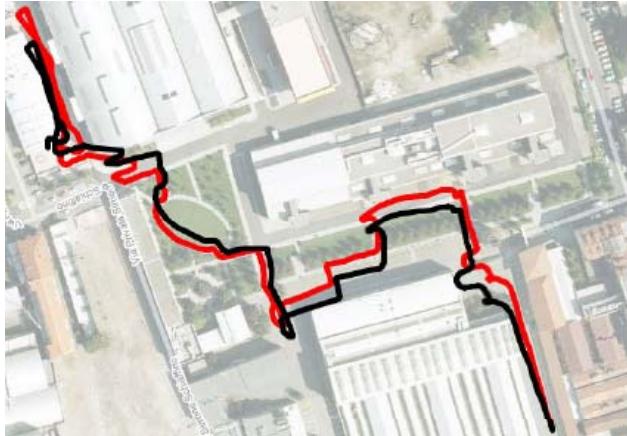


Visually impaired

Related Work

Simultaneous Localization and Mapping (SLAM)

- Build map and localize agent in the map at the same time



Visual odometry over 650m trajectory
(Civera, 2009)

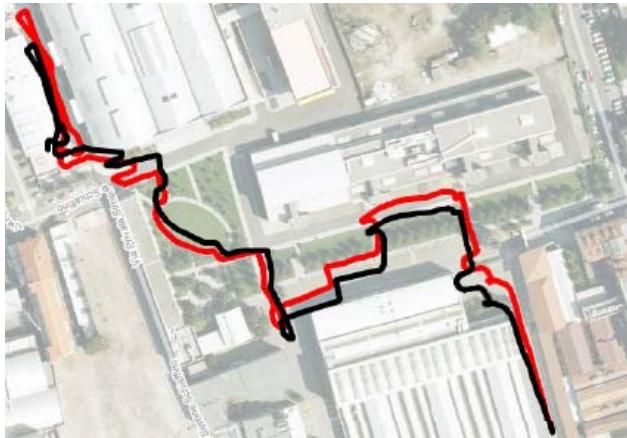


Laser-based SLAM, 30x30m (Grisetti, 2006)

Related Work

Simultaneous Localization and Mapping (SLAM)

- Build map and localize agent in the map at the same time



Visual odometry over 650m trajectory
(Civera, 2009)



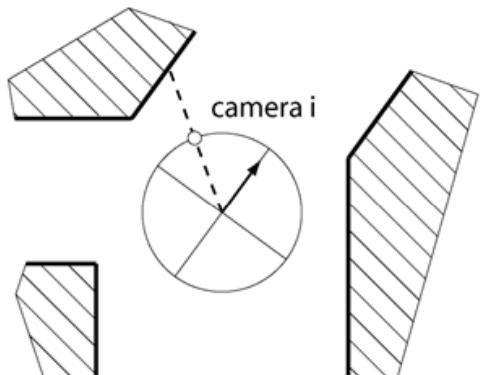
Laser-based SLAM, 30x30m (Grisetti, 2006)

Limitations of SLAM

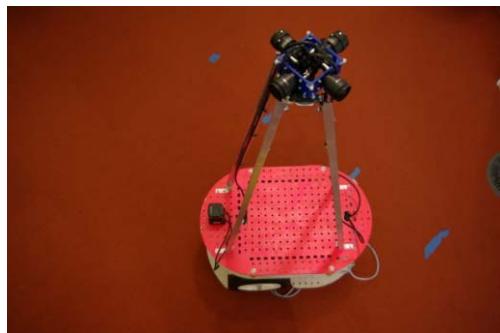
- Sensitive to degenerate user motions
- Sensitive to ambiguous world configurations
- Challenging in dynamic environments
- Requires full sensor calibration

Full metric reconstruction overkill for human navigation?

Contributions



- Sensor suite for vision-based navigation
- Body-relative navigation guidance using uncalibrated cameras
 - Topological mapping & loop closure
 - Localization & Rotation guidance
- Application to ground robot navigation
- User Study



System Overview



Four cameras, 752 x 480 8-bit images
Horizontal FOV: 360°



Sample frame, MIT Bldg 2

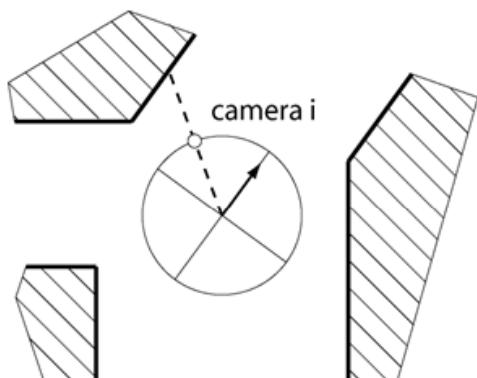
Why vision?

- Light, cheap, compact
- Rich information

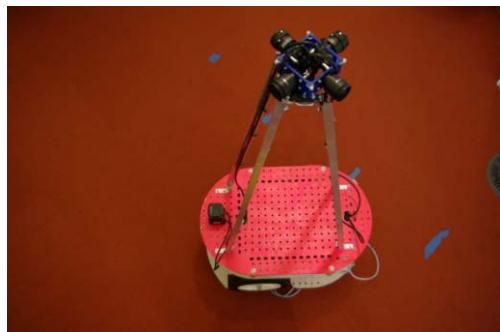
Why uncalibrated cameras?

- Intrinsic calibration is tedious
(*center of projection, focal length, distortion parameters*)
- Extrinsic calibration is hard
(*sensor-to-body transformation, 6DOF/camera*)
- Calibration subject to change in real conditions

Outline



- Sensor suite for vision-based navigation
- Body-relative navigation guidance using uncalibrated cameras
 - Topological mapping & loop closure
 - Localization & Rotation guidance
- Application to ground robot navigation
- User Study

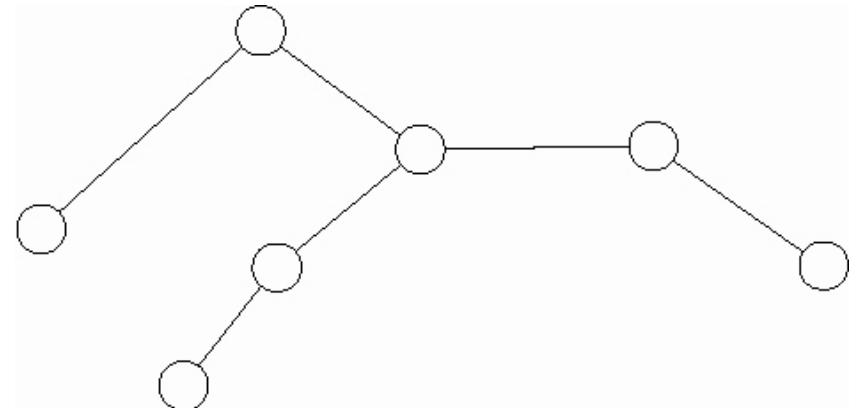


Method Overview

EXPLORATION (FIRST VISIT)

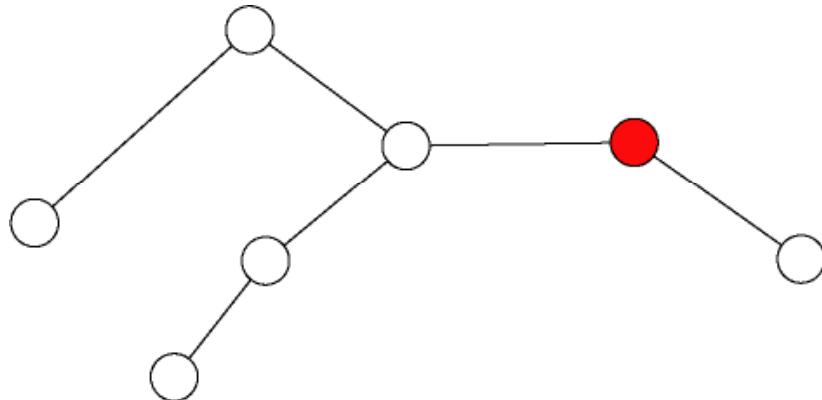
Place graph

- Node = location in the world
- Edge = physical path btw nodes
- Arbitrary 3D environments

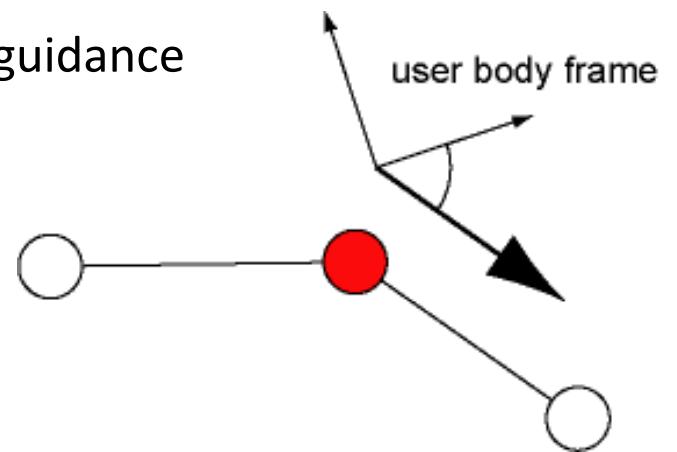


NAVIGATION (SUCCESSIVE VISIT)

Localization (node estimation)

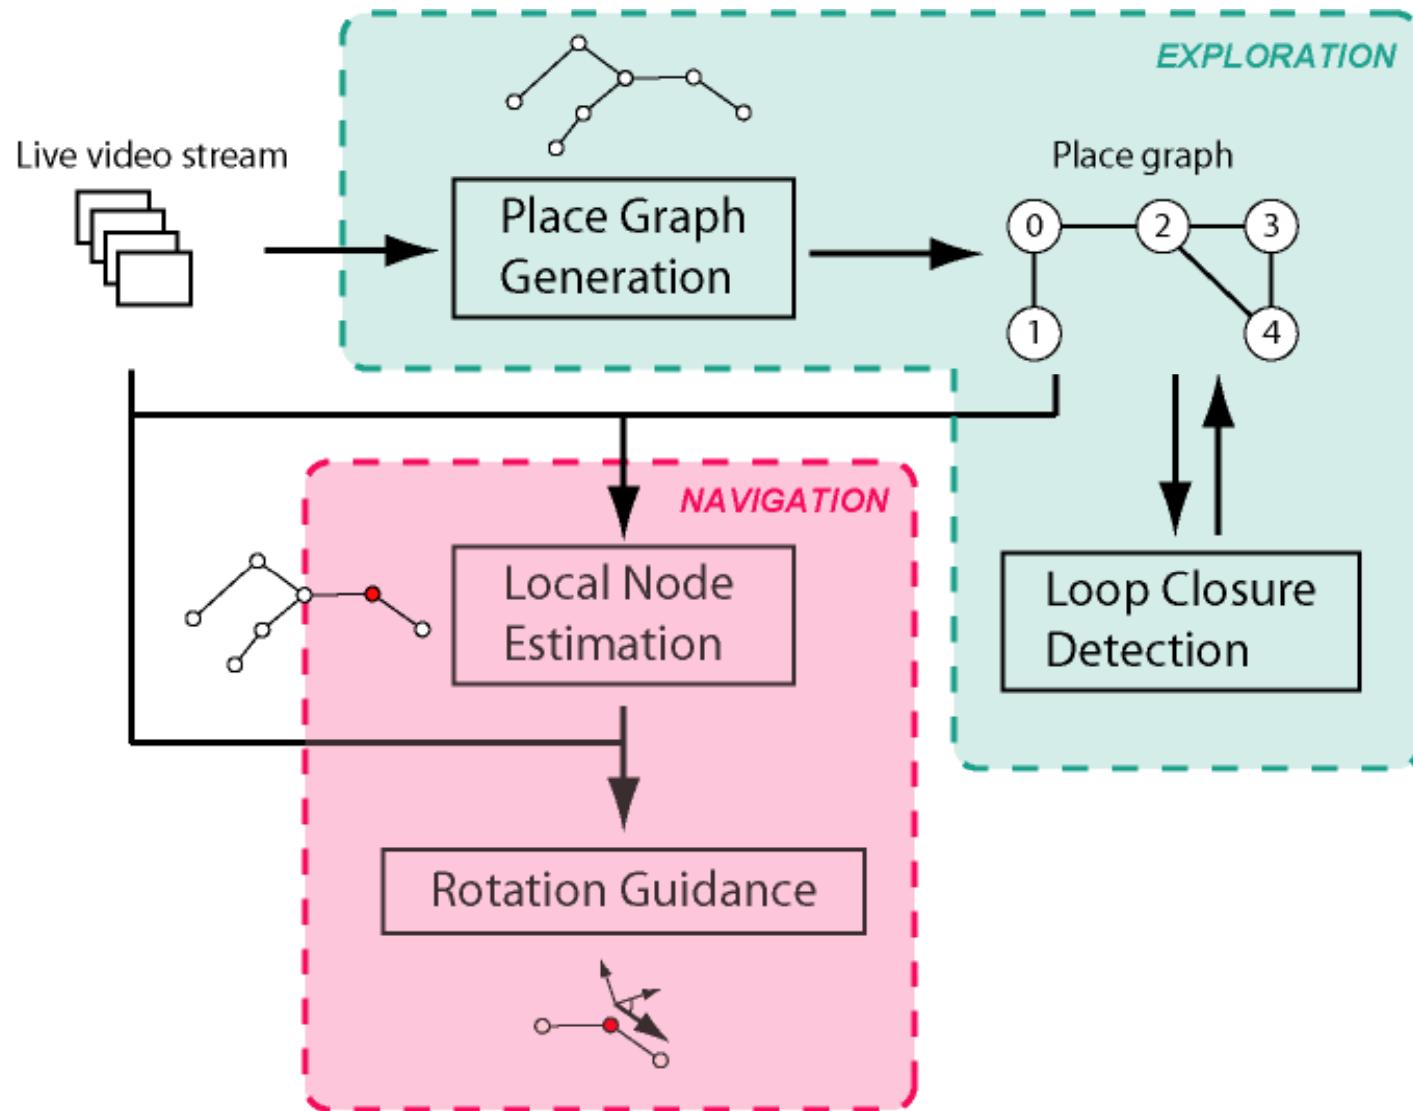


Rotation guidance



Body-relative navigation guidance using uncalibrated cameras, O. Koch, S. Teller, ICCV 2009

Method Overview

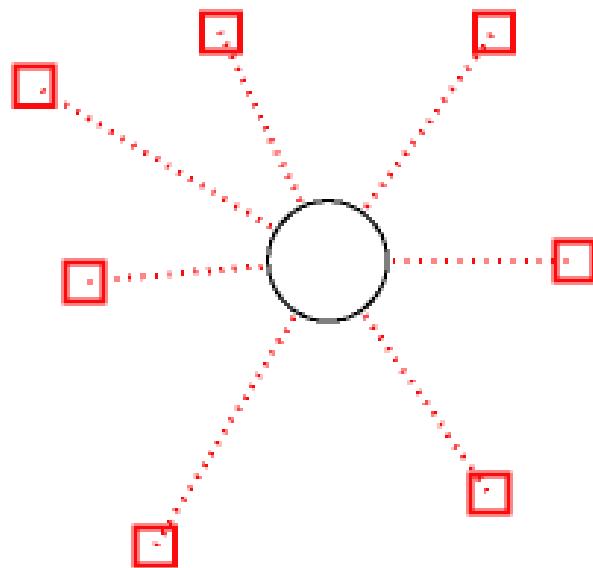


Body-relative navigation guidance using uncalibrated cameras, O. Koch, S. Teller, ICCV 2009

Place Graph

Place Graph
Generation

- Sparse representation of the exploration path



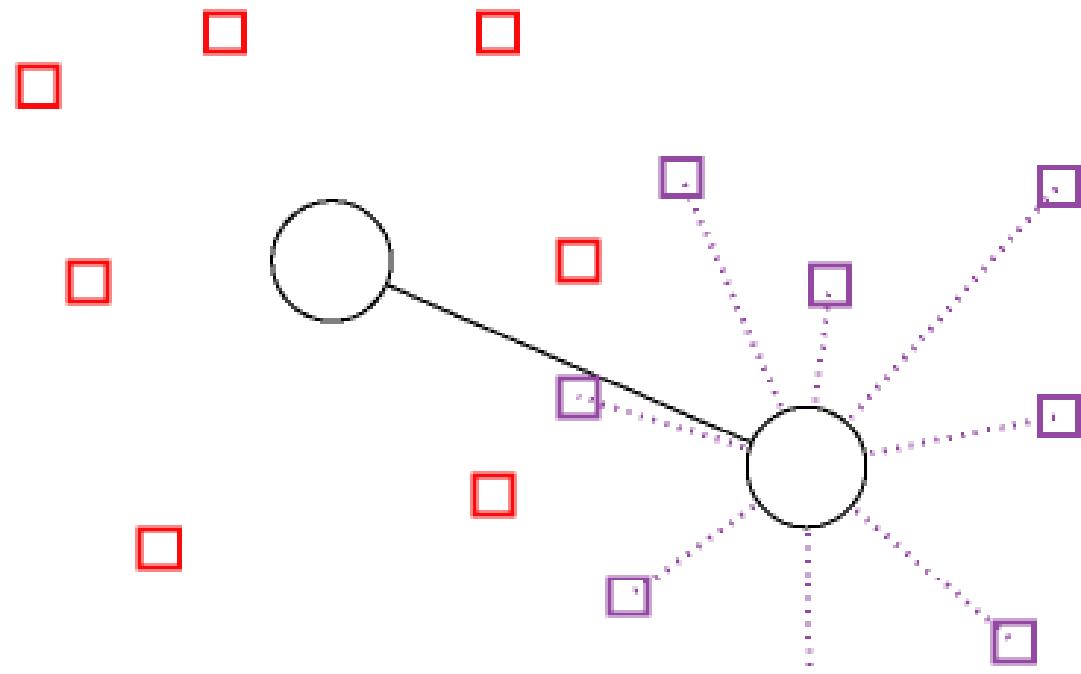
□ World features

○ Graph node

Place Graph

Place Graph
Generation

- Sparse representation of the exploration path



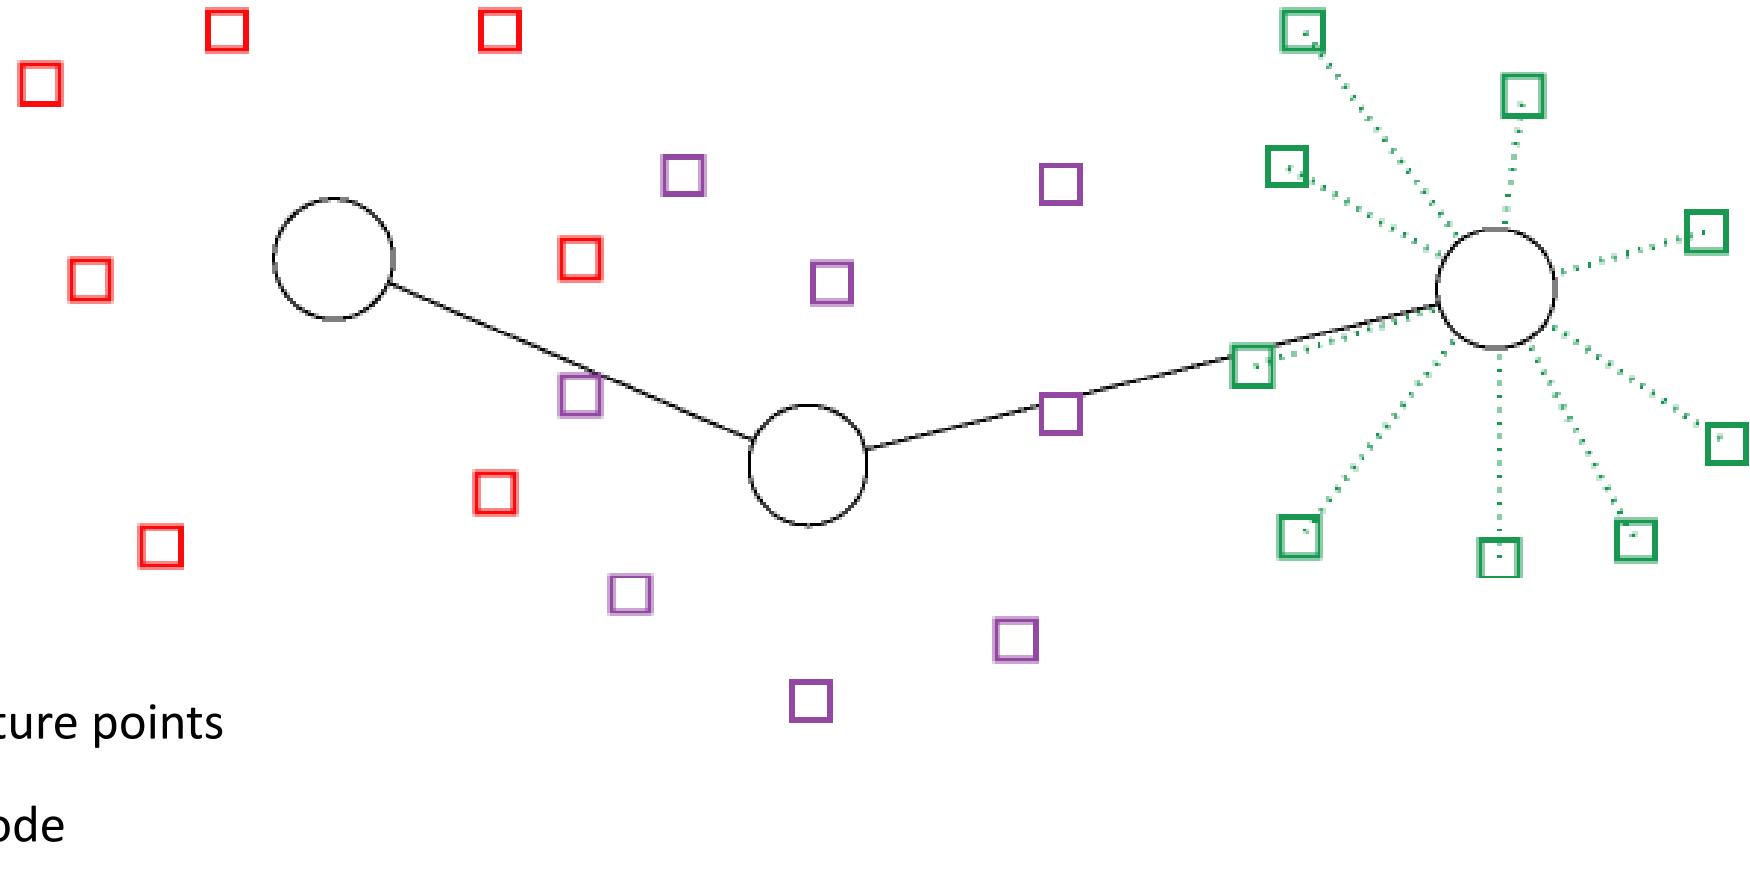
World features

Node

Place Graph

Place Graph
Generation

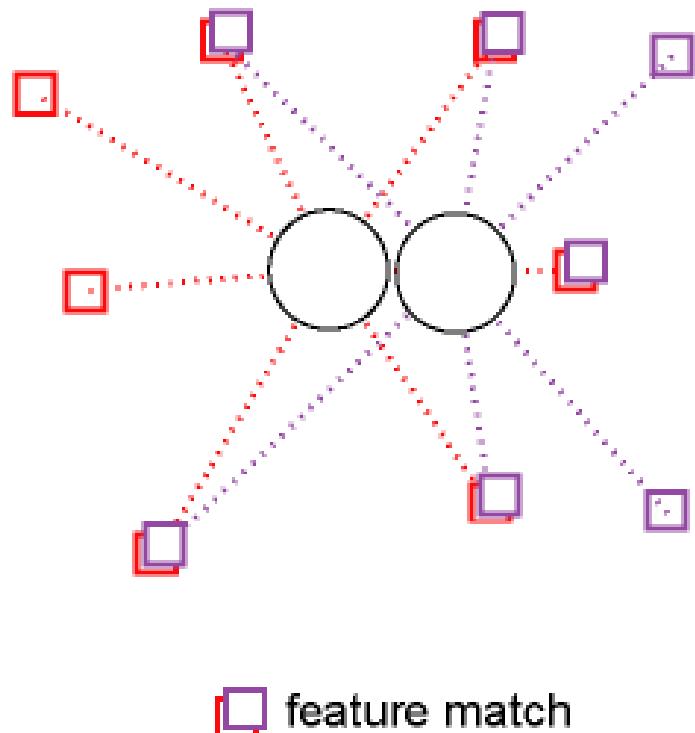
- Sparse representation of the exploration path



Place Graph Generation

Place Graph
Generation

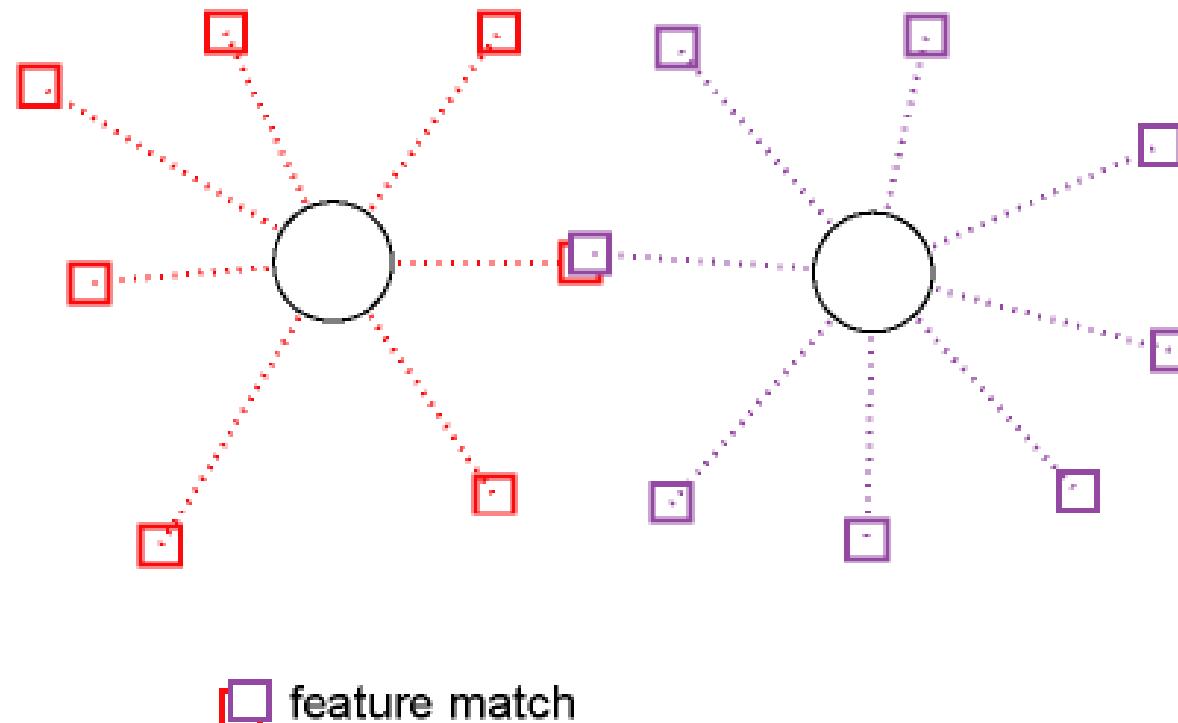
- Variation in environment appearance is captured by feature matching



Place Graph Generation

Place Graph
Generation

- Variation in environment appearance is captured by feature matching



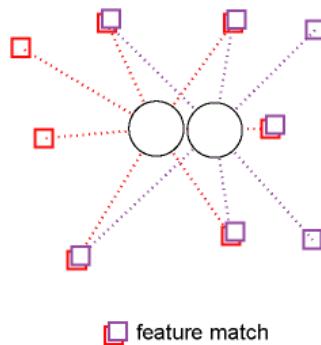
Place Graph Generation

- Energy function Ψ measures visual similarity.

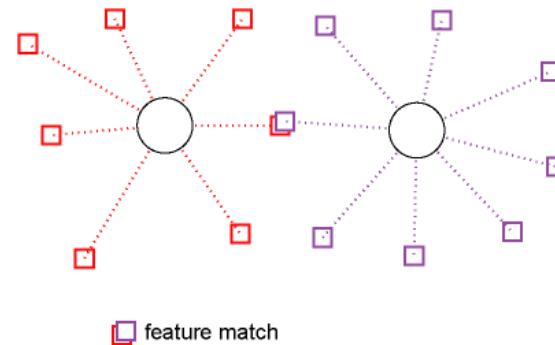
$$\Psi(F_1, F_2) = \frac{1}{|F_1| + |F_2|} \cdot \left(|F_1| + |F_2| - 2 \cdot |\Phi(F_1, F_2)| + \sum_{(f_1^i, f_2^j) \in \Phi(F_1, F_2)} \|f_1^i - f_2^j\|_2 \right)$$

F1, F2: feature sets
penalty for unmatched features
average L2 distance in feature space

Φ : feature matching function



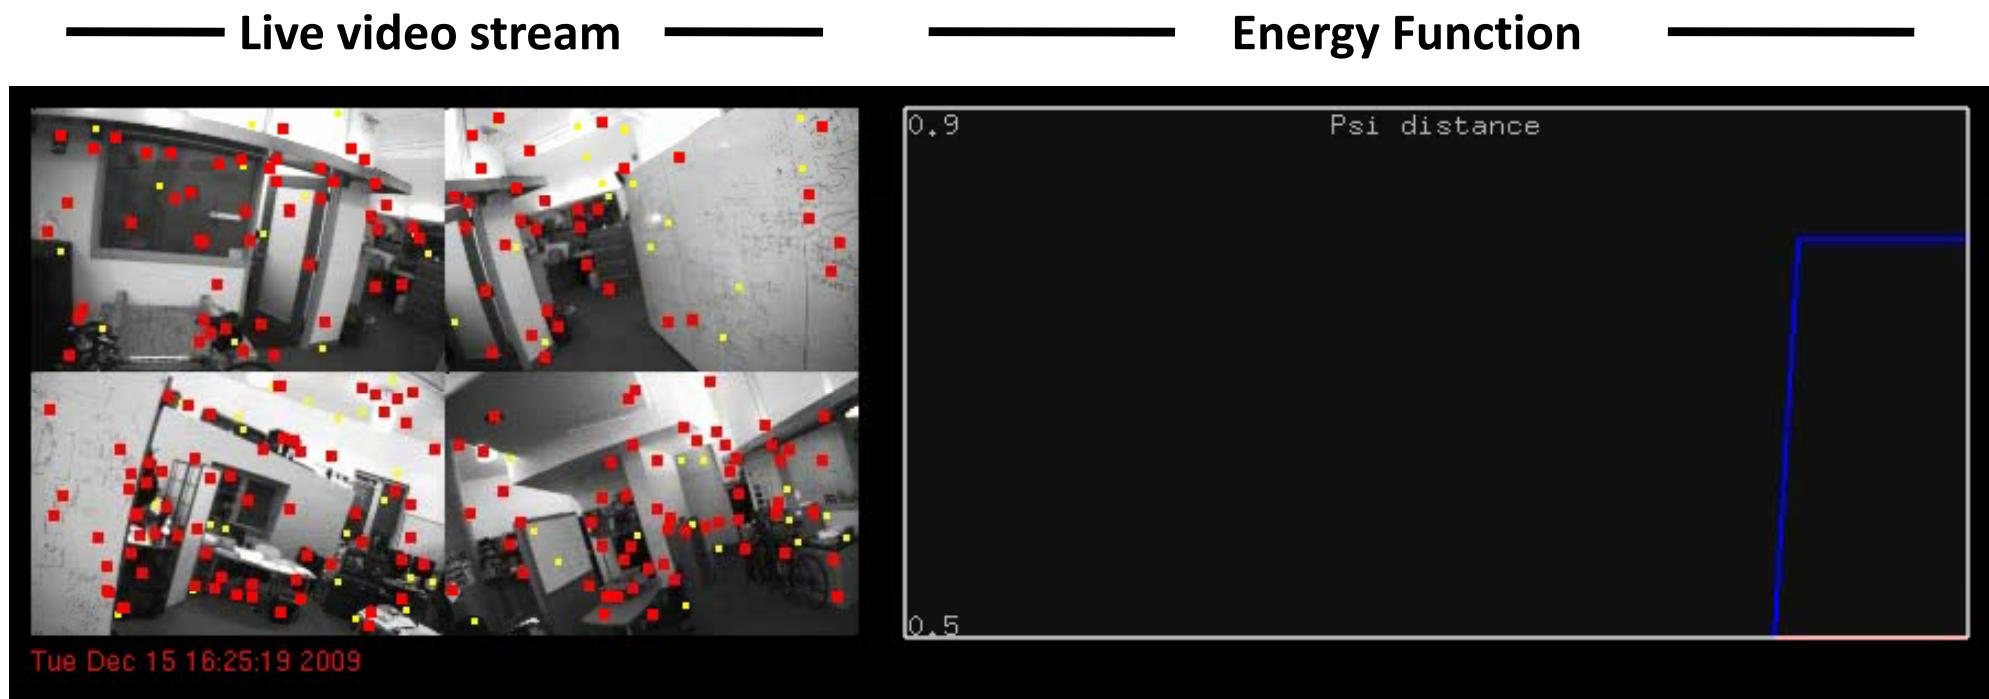
$$0 < \Psi < \Psi_0$$



$$\Psi_0 < \Psi < 1$$

Place Graph Generation

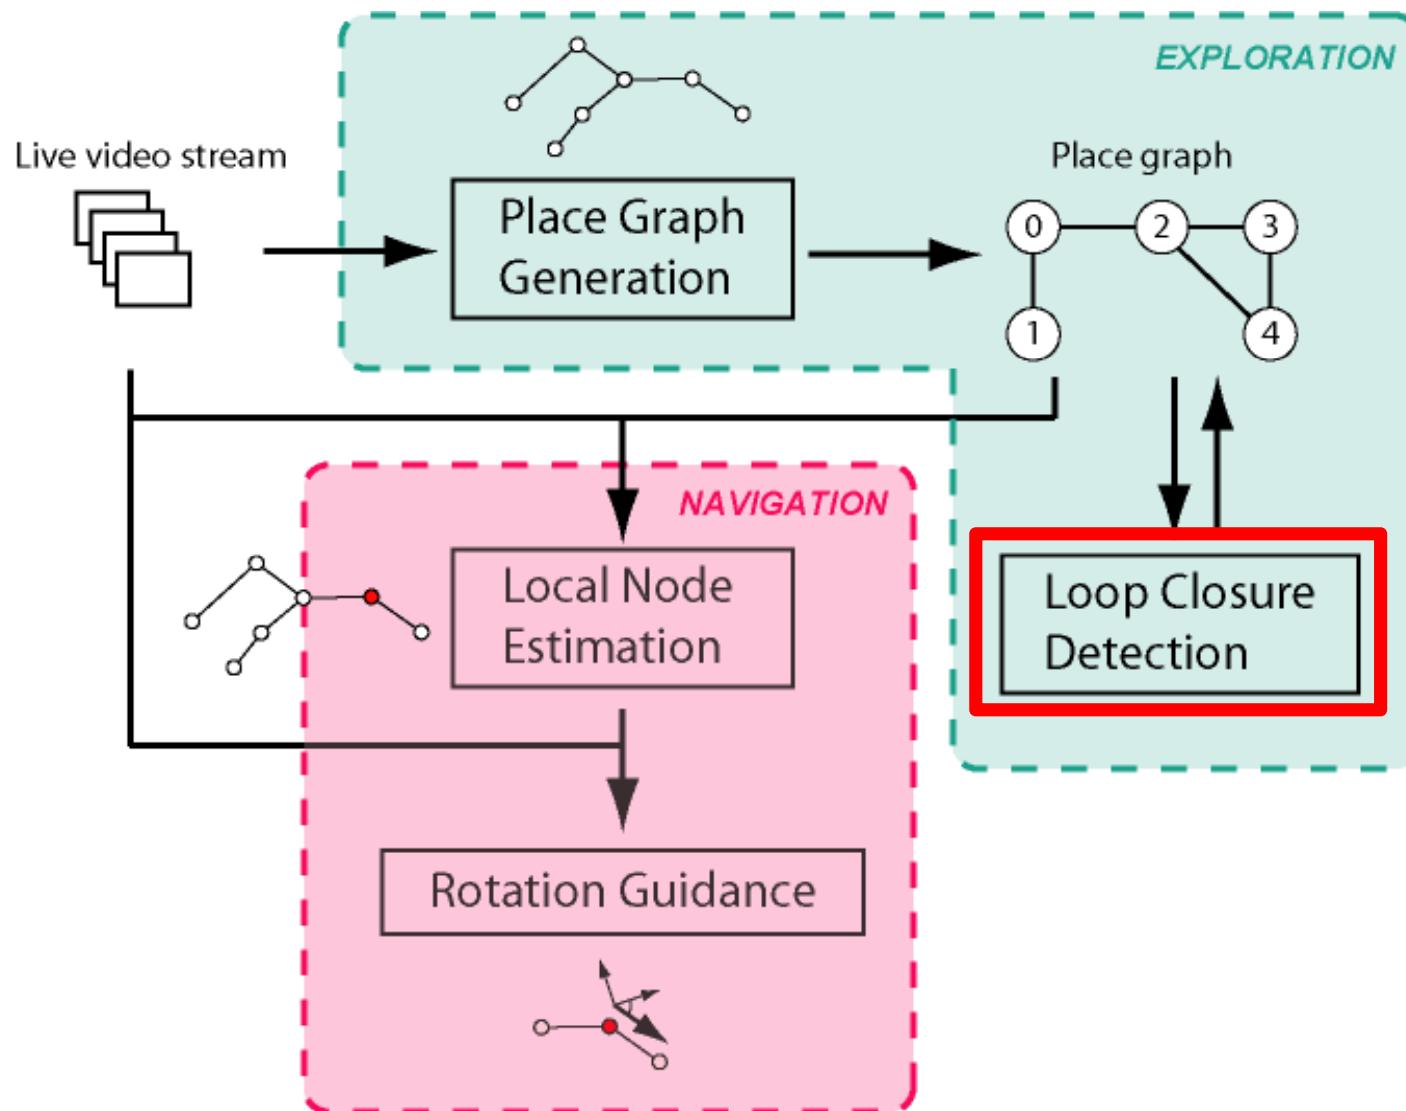
Place Graph
Generation



■ Features

■ Matched features

Method Overview

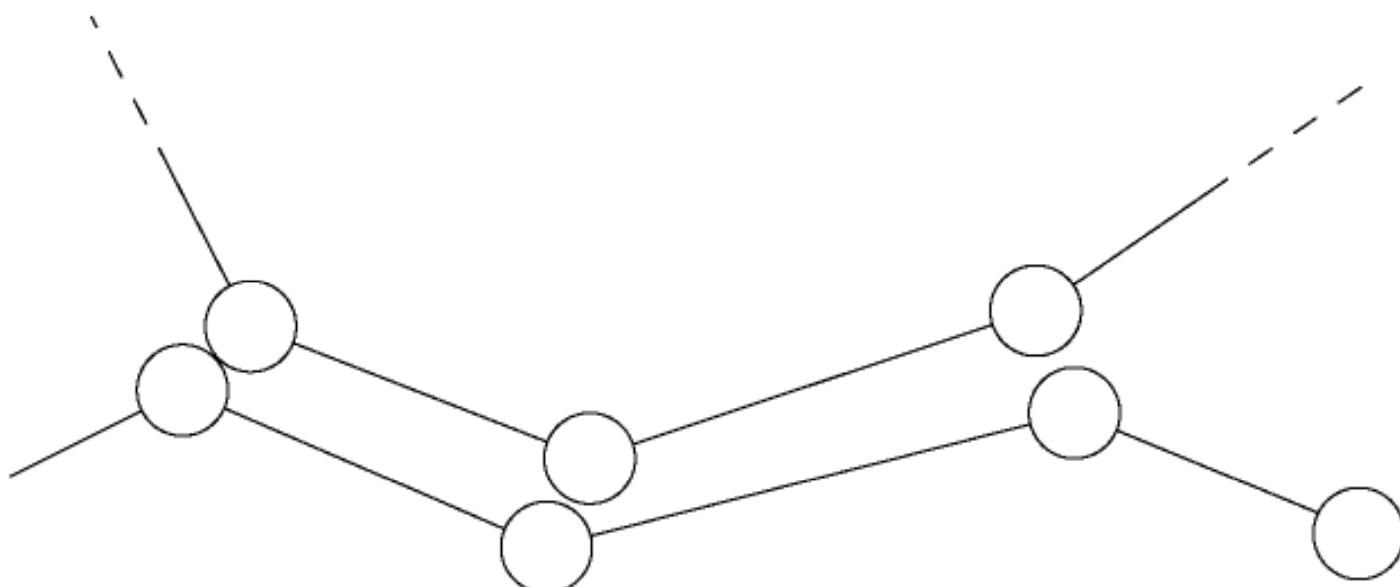


Loop Closure Detection

Loop Closure Detection

Approach

1. generate loop closure hypothesis using visual similarity between nodes
 2. validate hypothesis using temporal consistency



Bag-of-words

- Bag-of-words: model for data representation

- John will fly to San Francisco next week.
- Mary would like to fly to the moon with John.

Dictionary: {1: John, 2: will, 3: fly, 4: like, 5: to, 6: moon, 7: Francisco, 8: week, 9: Mary}

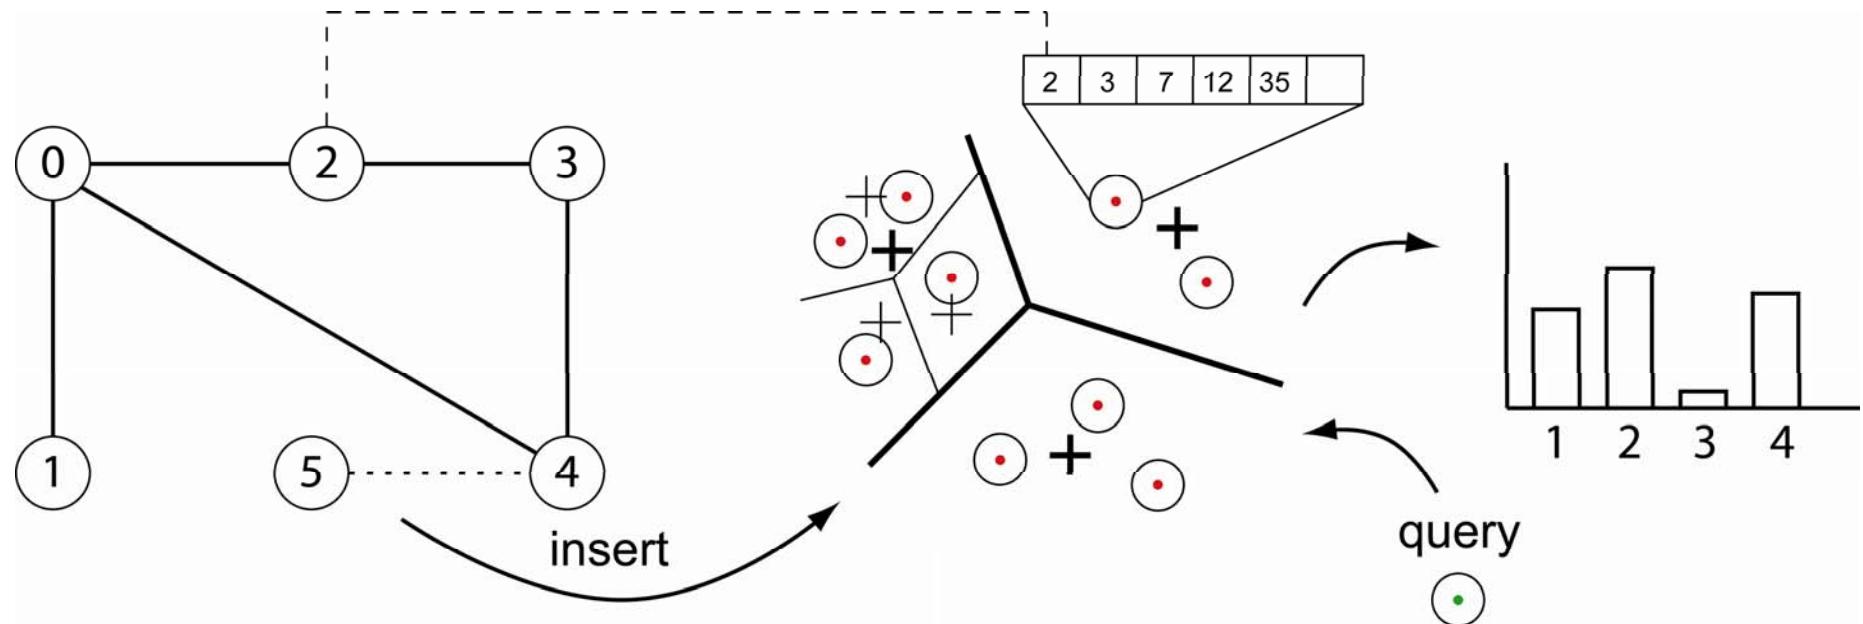
- [1, 1, 1, 0, 1, 0, 1, 1, 0]
- [1, 0, 1, 1, 1, 1, 0, 0, 1]

- Image=“document”, image feature=“word”
- Applications
 - object categorization (generative/discriminative models)
 - object segmentation, localization

Loop Closure Detection

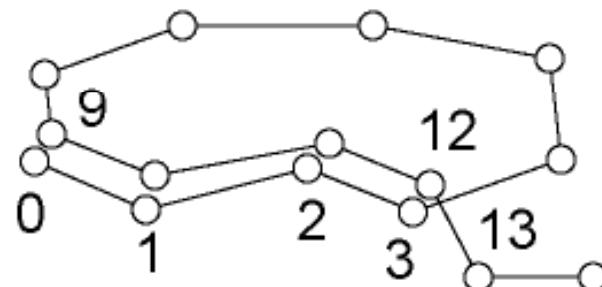
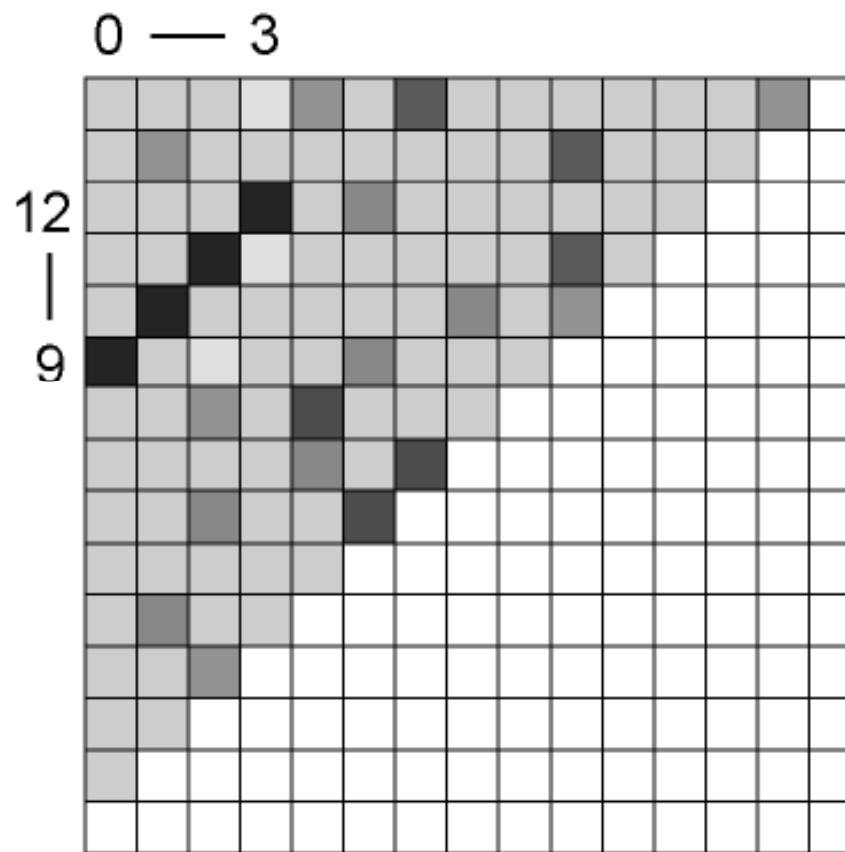
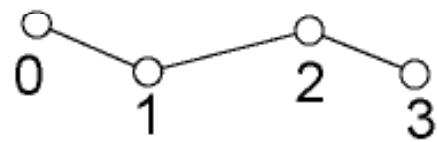
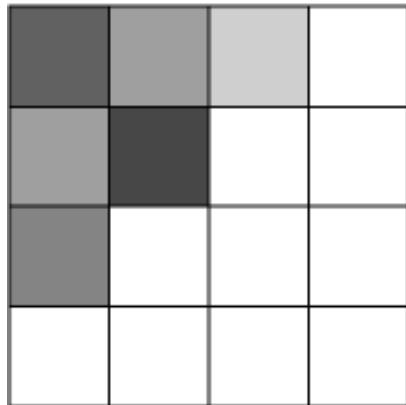
Loop Closure
Detection

- Node=“document”, image feature=“word”



Loop Closure Detection

- Similarity matrix



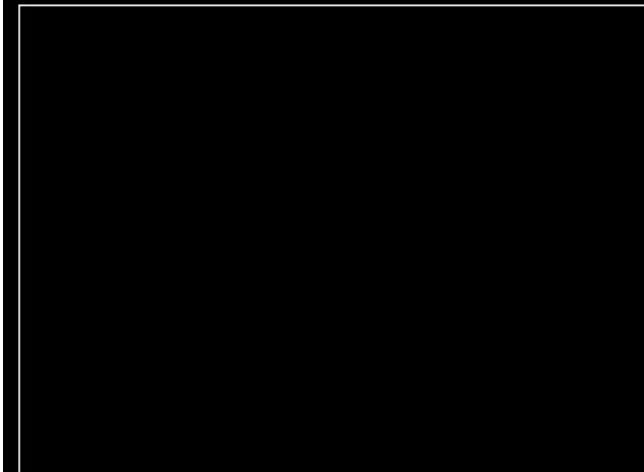
Loop Closure Detection

Loop Closure
Detection

Live video stream



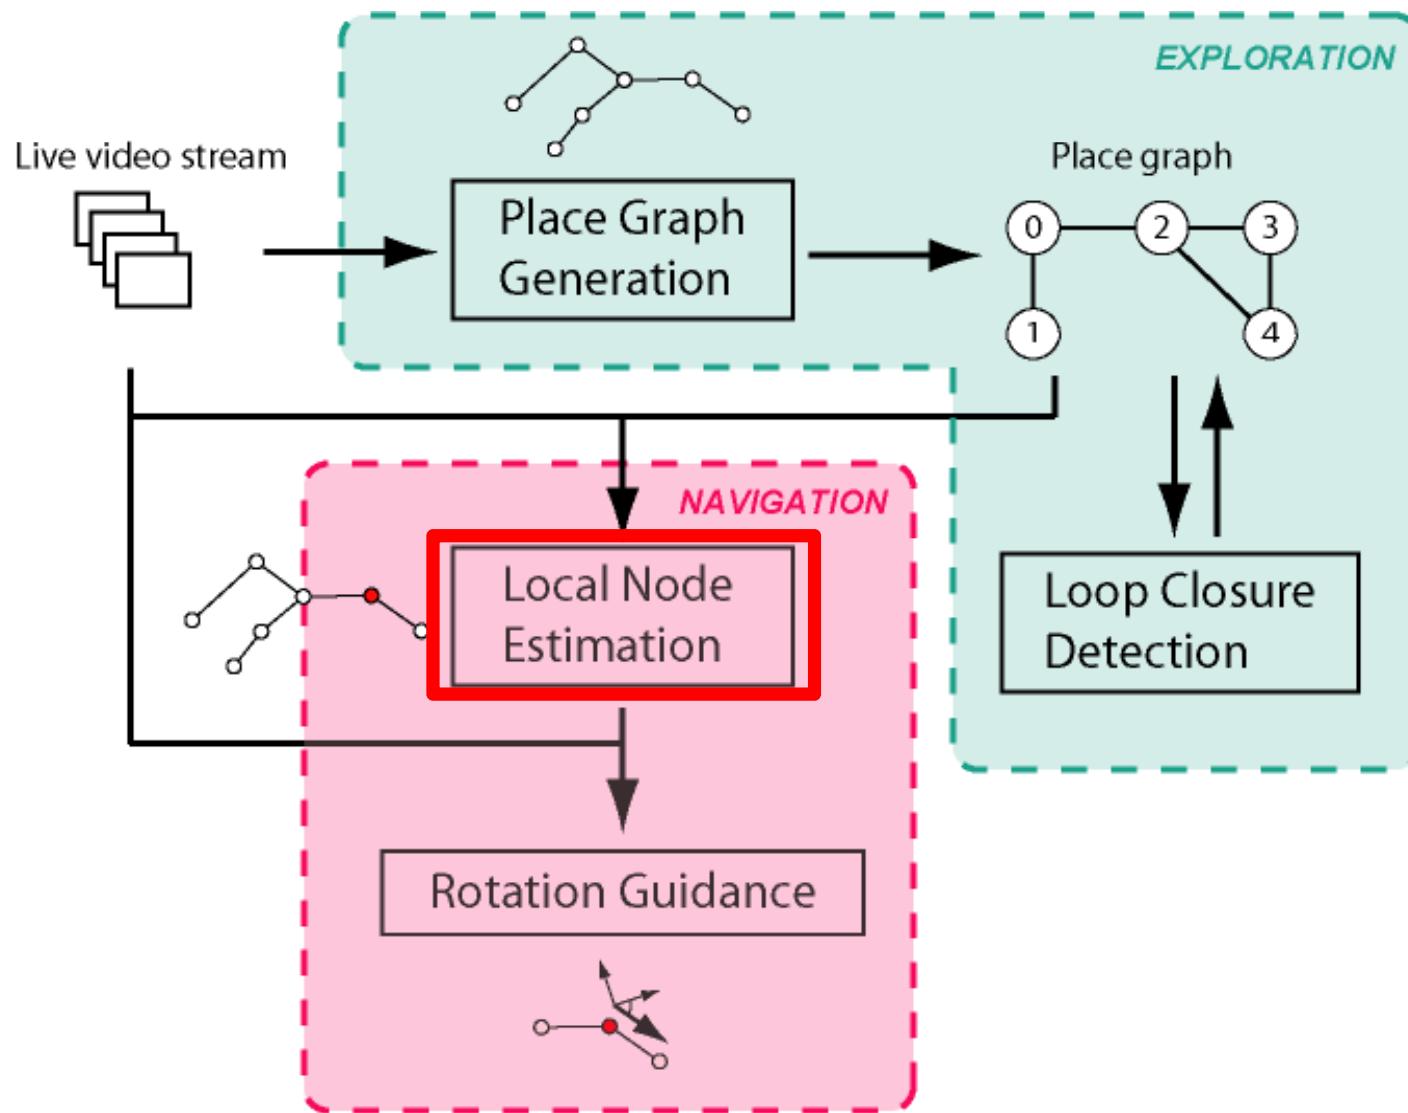
Wed Dec 2 12:40:45 2009



Place Graph

Similarity Matrix

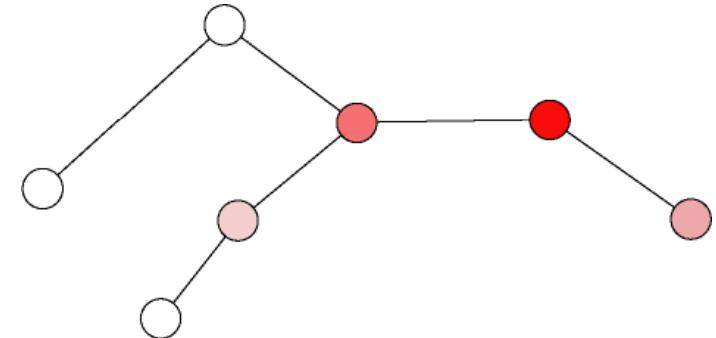
Method Overview



Local Node Estimation

Discrete recursive Bayesian filtering

$p(x_k \mid z_k)$: probability of state x_k
while observing z_k at time k



1. *Prediction step*

$$p(x_{k+1} \mid z_k) = \sum_k p(x_{k+1} \mid x_k) p(x_k \mid z_k)$$

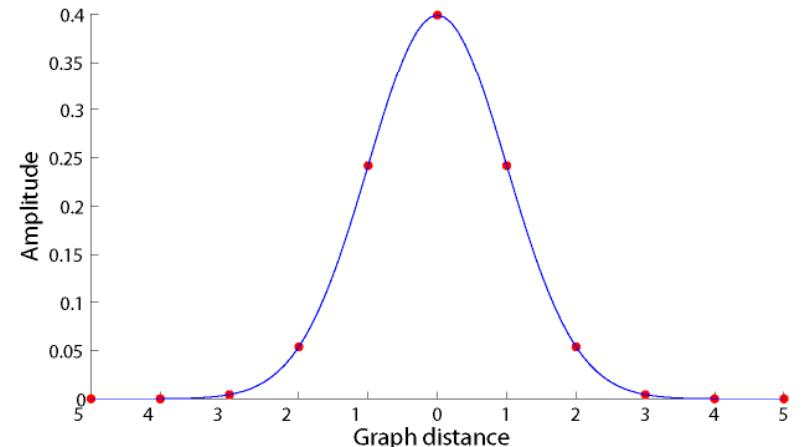
2. *Observation step*

$$p(x_{k+1} \mid z_{k+1}) = \lambda p(z_{k+1} \mid x_{k+1}) p(x_{k+1} \mid z_k)$$

Local Node Estimation

1. *Prediction step*

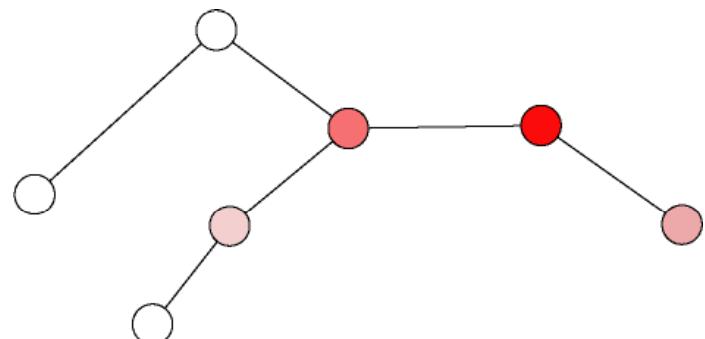
$p(x_{k+1} \mid x_k)$: motion model



Gauss Window Function

2. *Observation step*

$p(z_{k+1} \mid x_{k+1})$: observation model



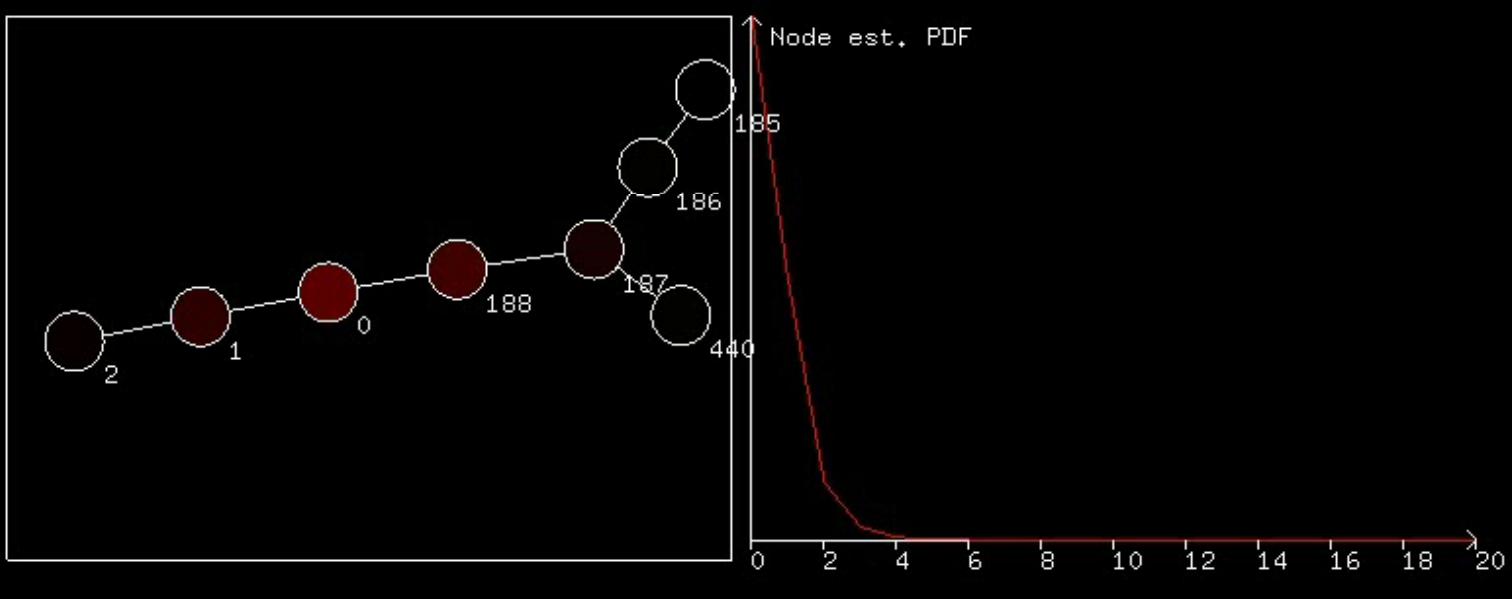
Local Node Estimation

Live video stream



Thu Mar 5 13:49:47 2009

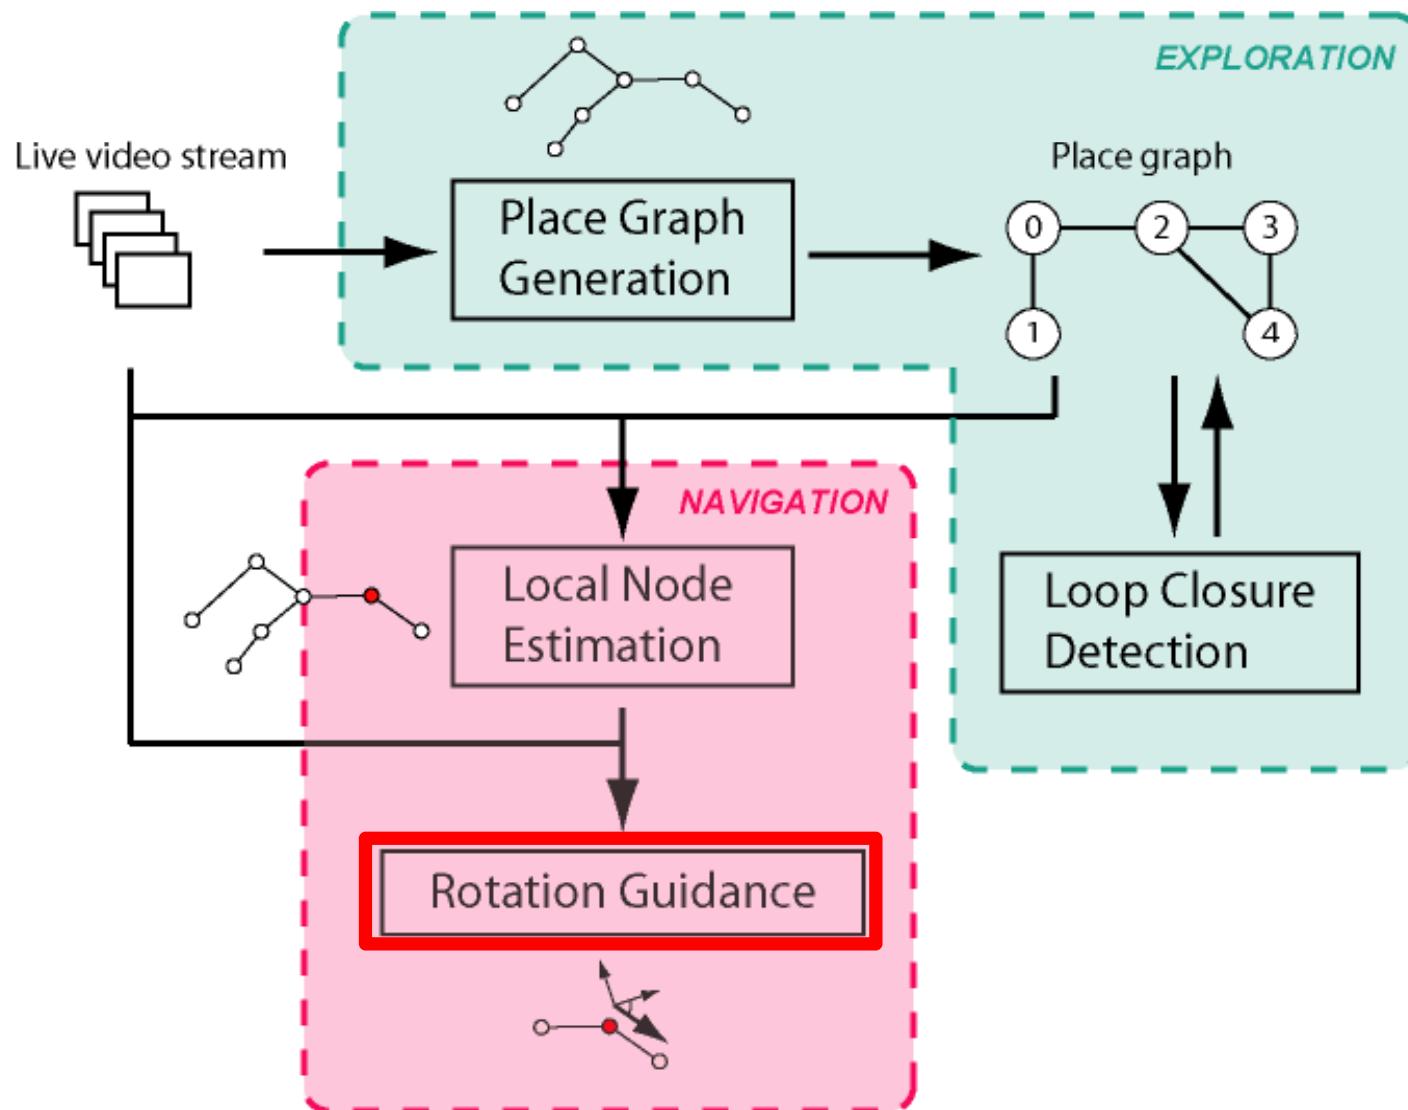
Localization



Place Graph

Node Estimation PDF

Method Overview

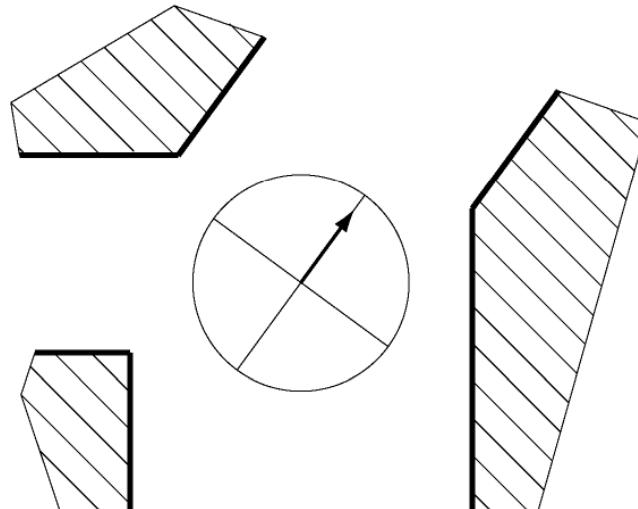


Rotation Guidance

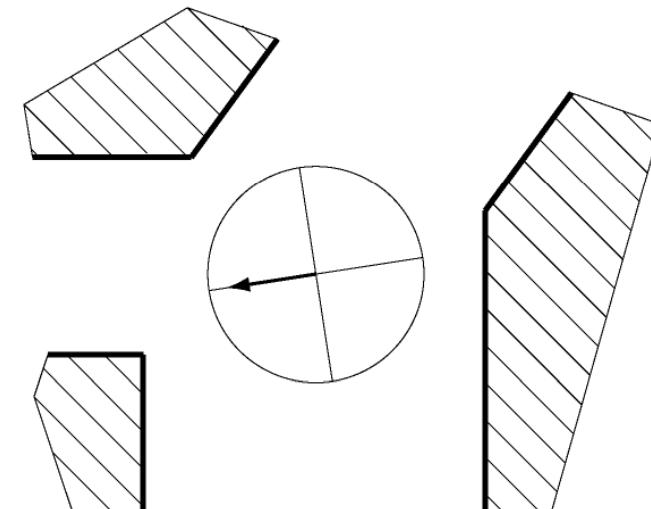
Relative orientation problem

Determine the orientation of the user at time t' *relative to* the orientation at time t .

time t



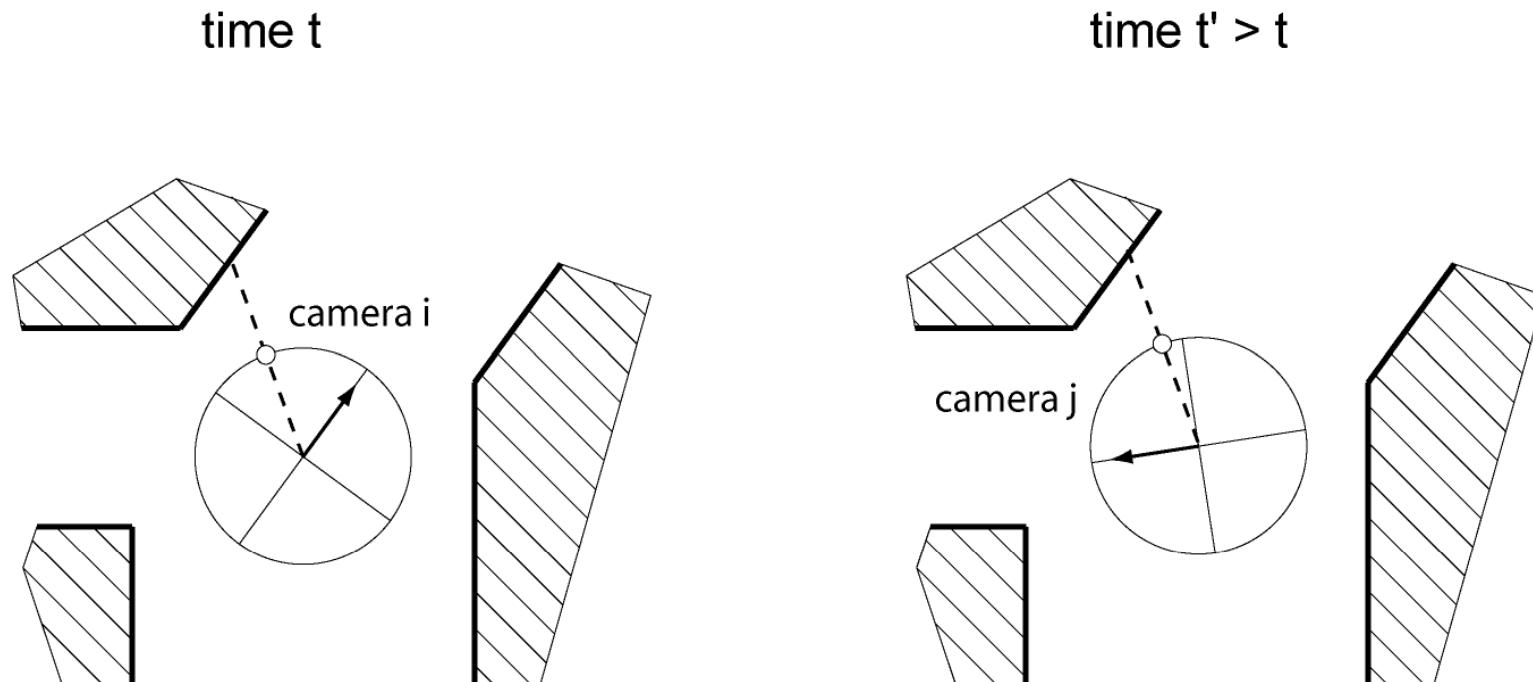
time $t' > t$



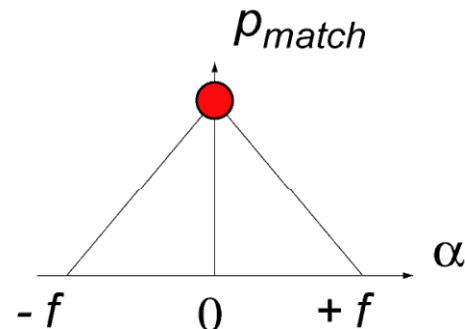
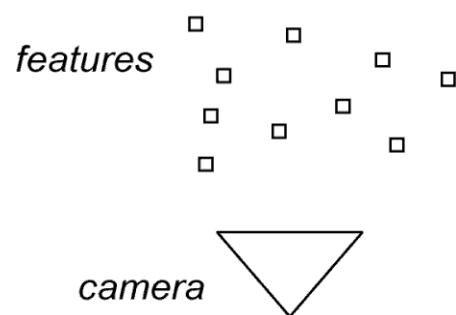
Rotation Guidance

Idea

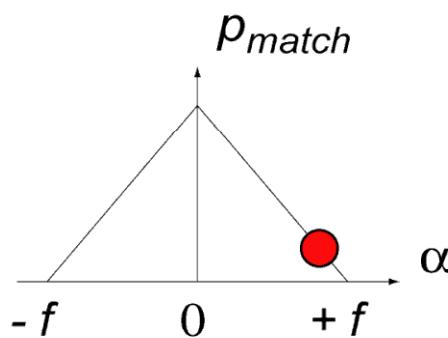
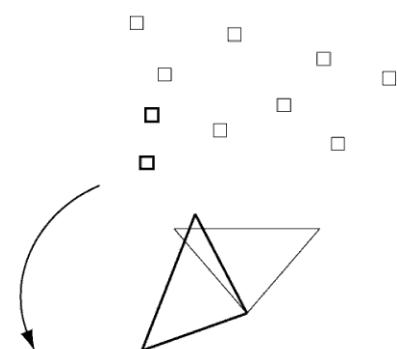
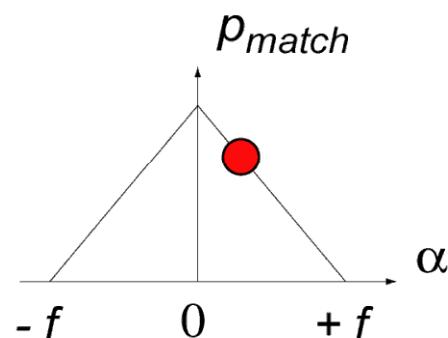
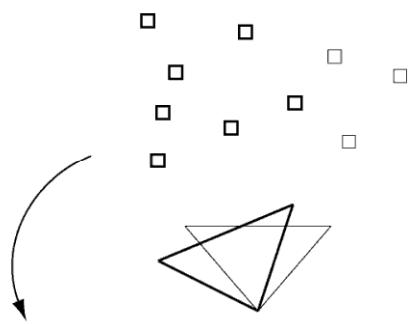
Learn the relationship between user rotation and feature matches across cameras.



Feature matching with one camera

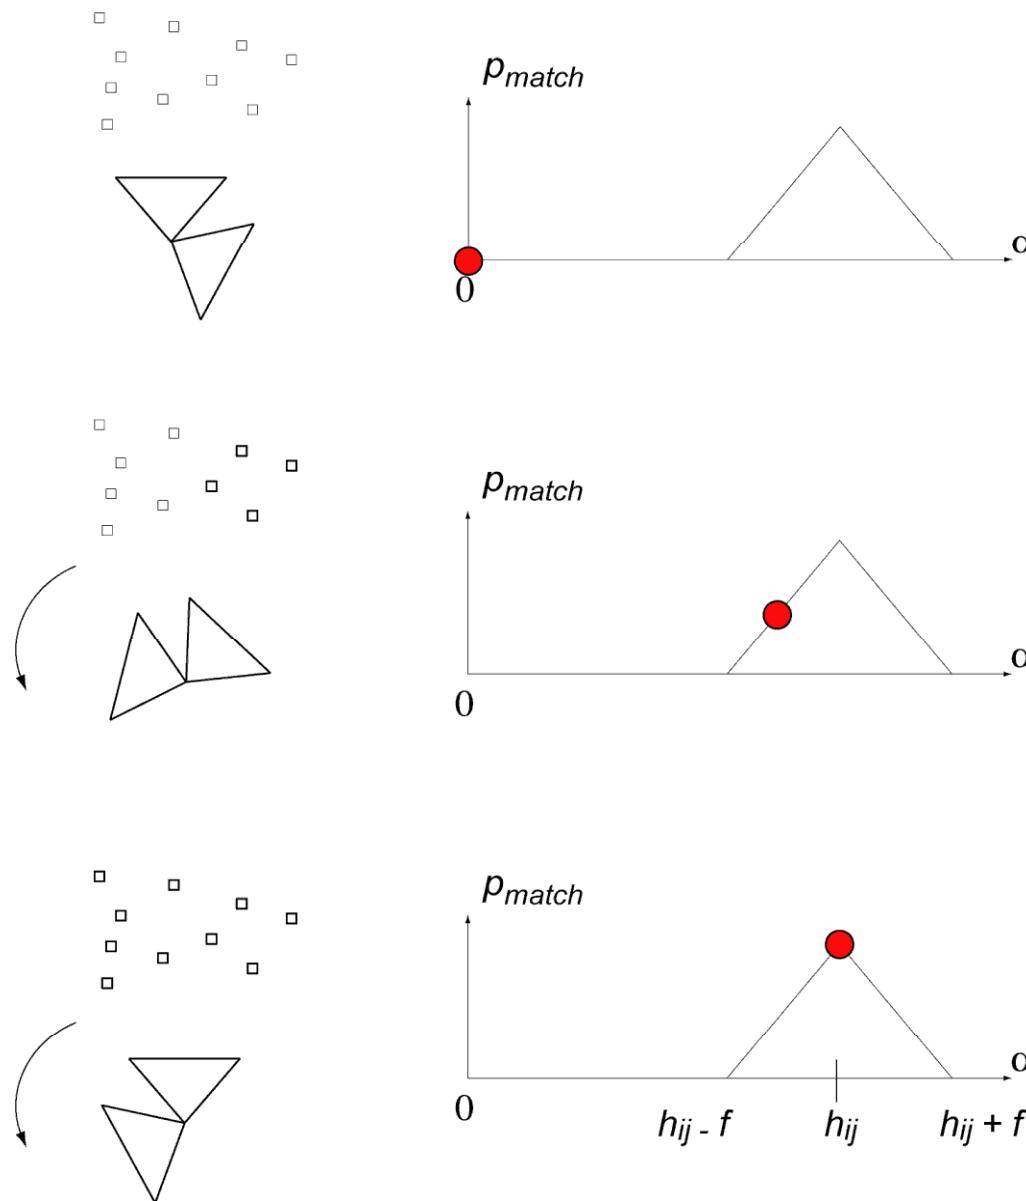


α : rotation angle
 f : camera field of view



Assumption: features are uniformly distributed in image space.

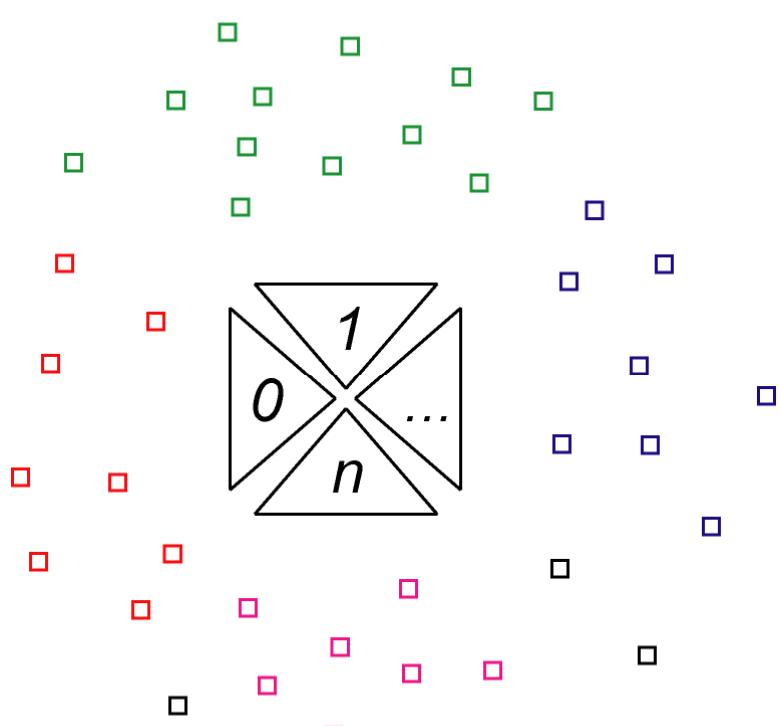
Feature matching with two cameras



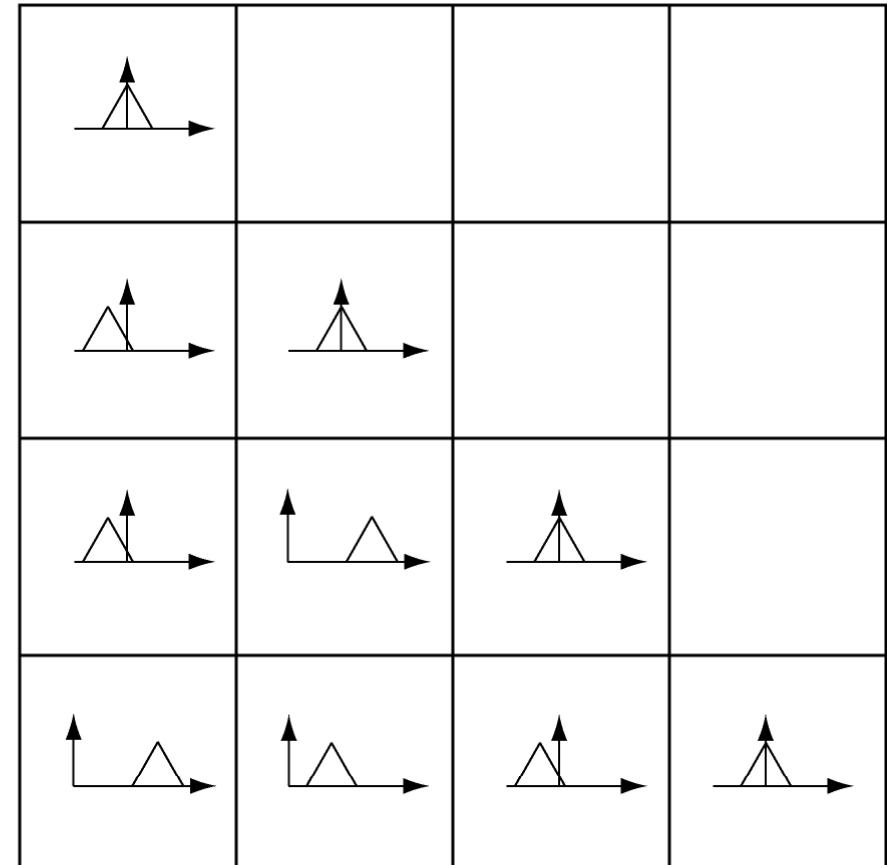
Assumption: cameras have the same field-of-view.

Feature matching with n cameras

Rotation
Guidance

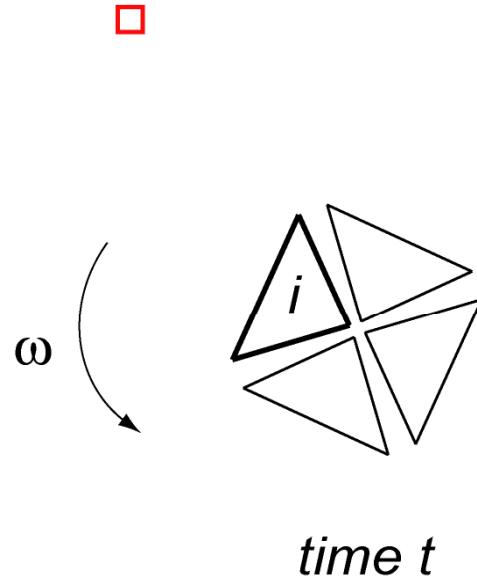


Match matrix ($n \times n$)



Learning the match matrix from training

System rotating in place at constant speed ω

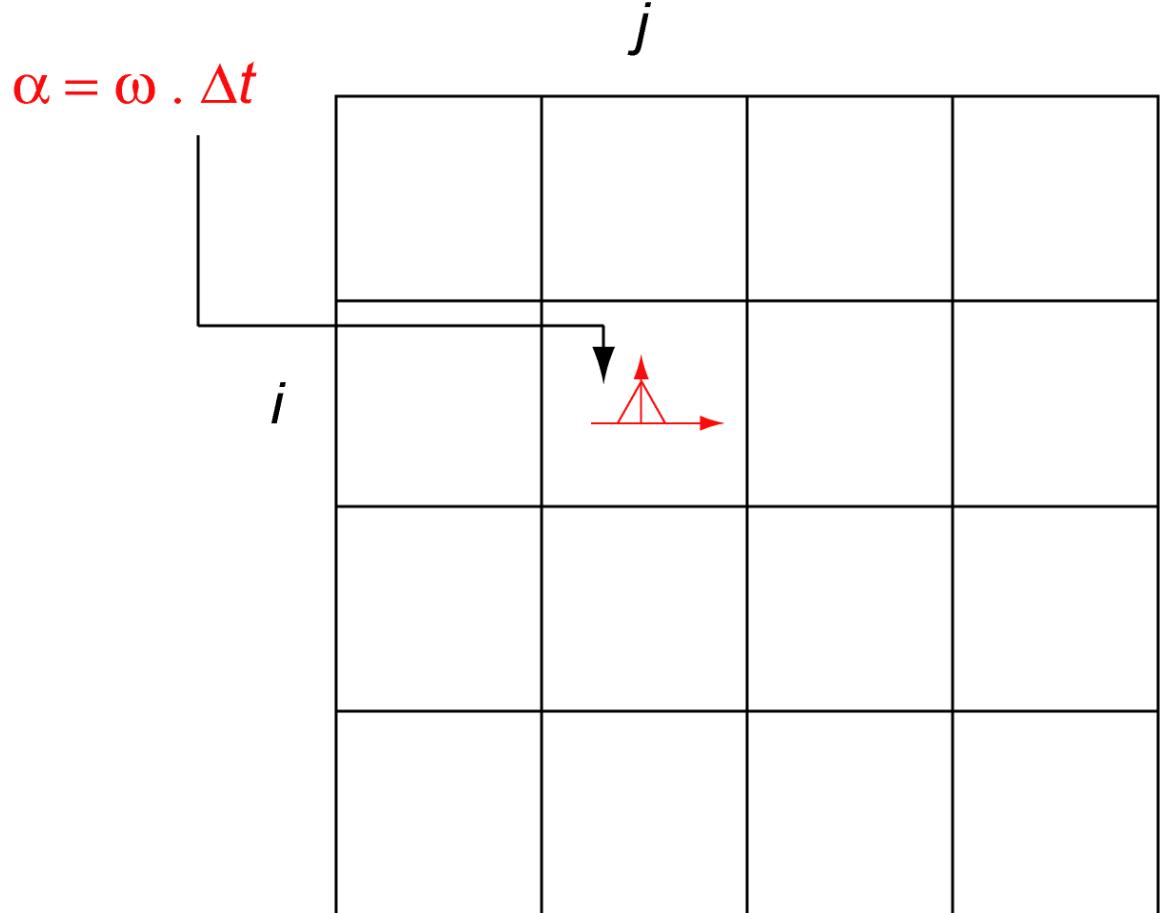
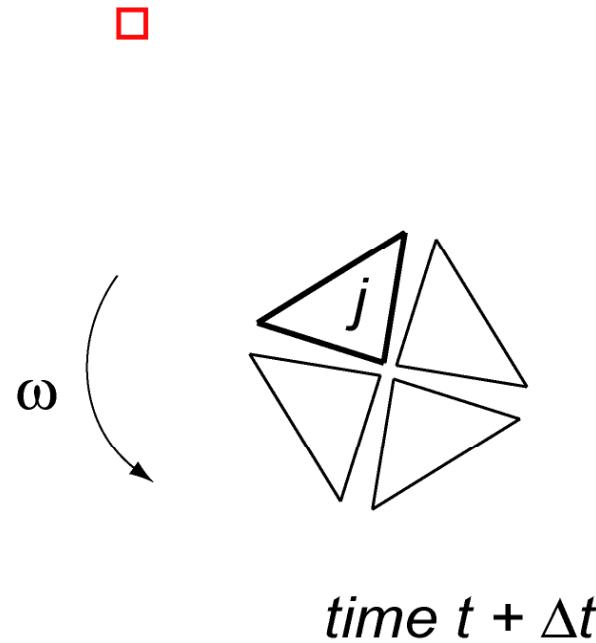


Match matrix

| | | | |
|--|--|--|--|
| | | | |
| | | | |
| | | | |
| | | | |

- *feature observed on camera i at time t*

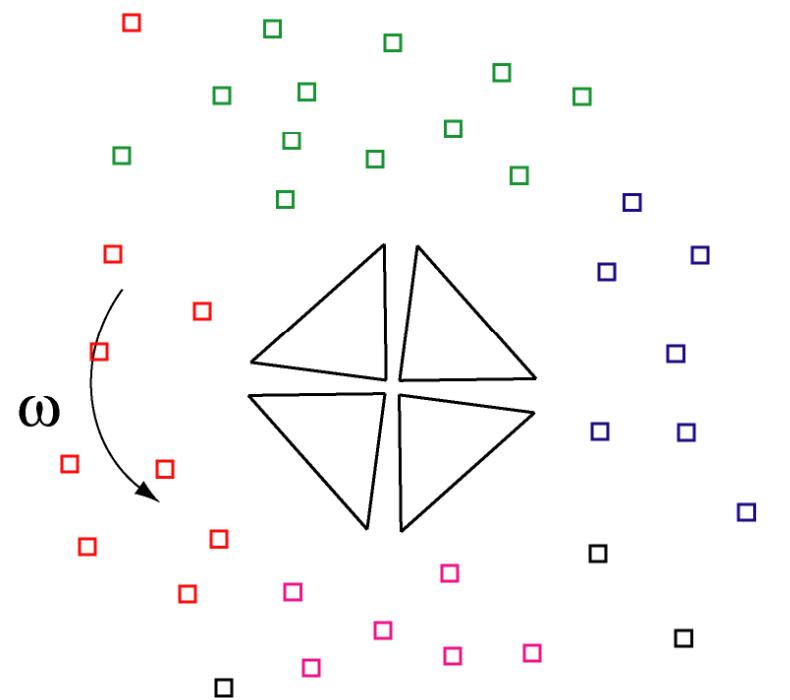
Learning the match matrix from training



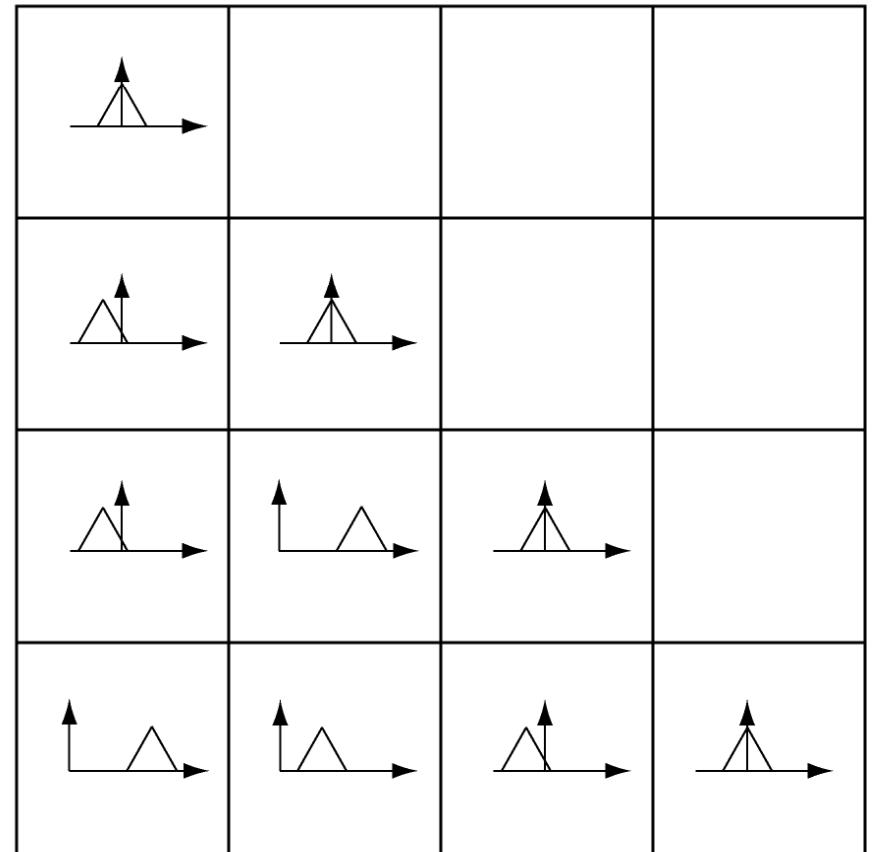
- (same) feature observed on camera j at time $t + \Delta t$

Learning the match matrix from training

Repeat for all features and all pairs of image frames (time sampling)



Match matrix ($n \times n$)



Rotation Guidance

Match Matrix

Live video stream

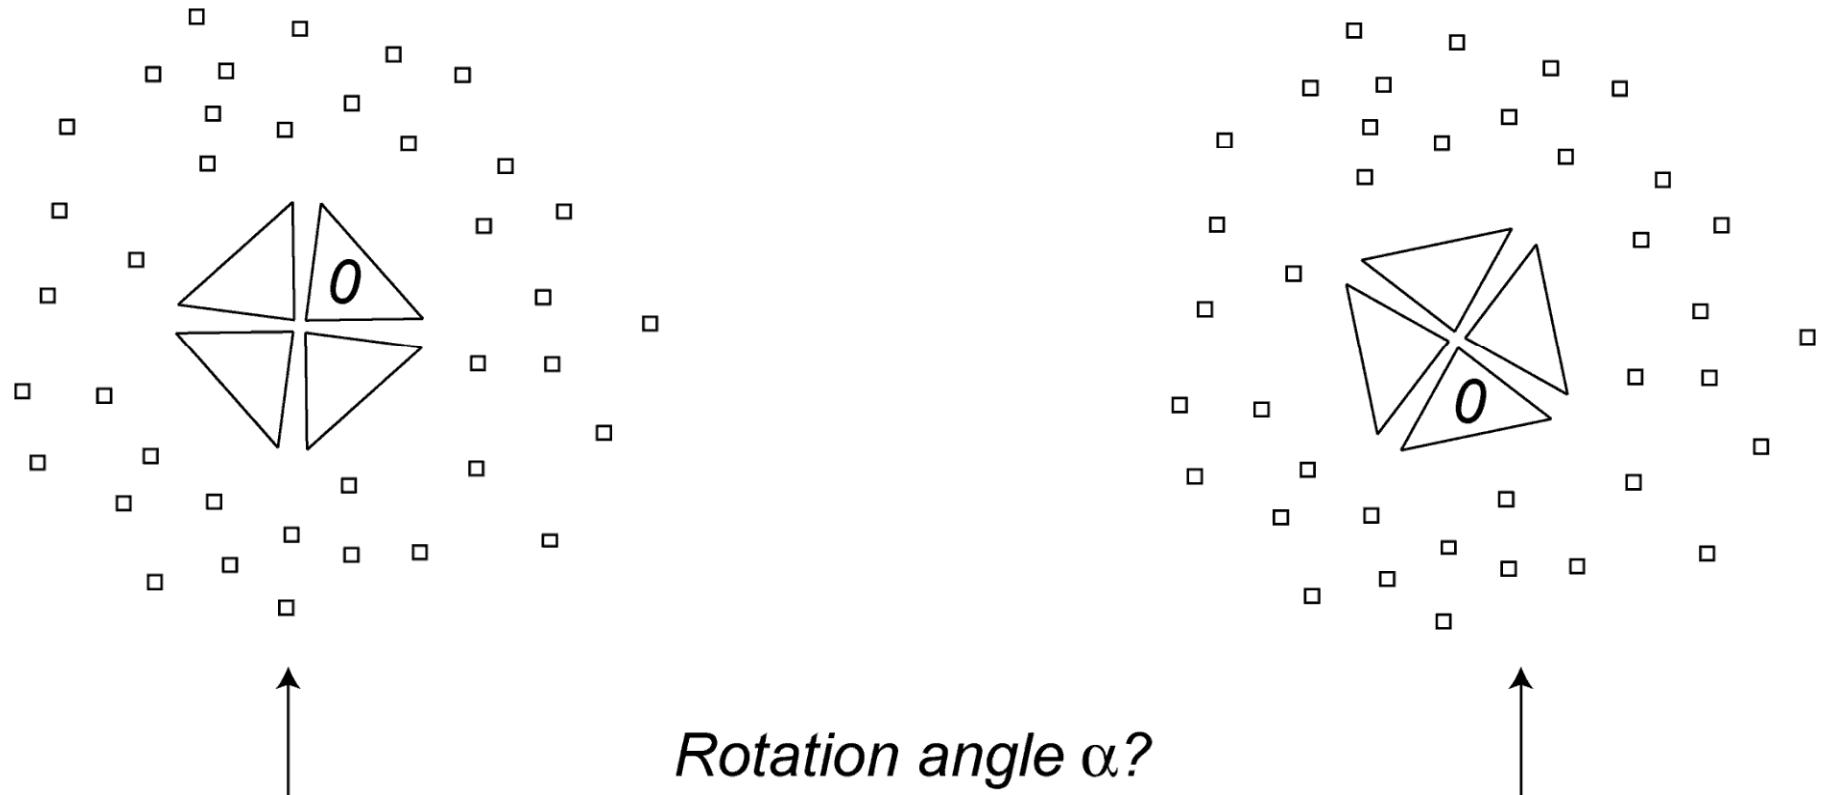


Close-up view of matrix cell

Relative orientation problem

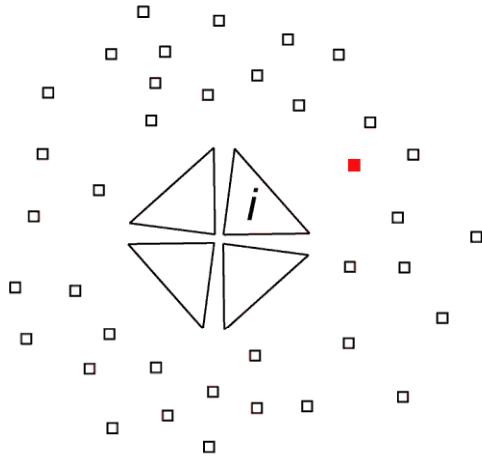
Relative orientation problem

Determine the orientation of the user at time t' *relative to* the orientation at time t.

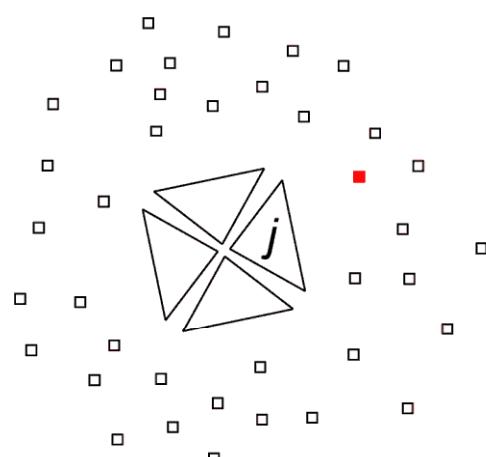


Orientation from a single correspondence

time t



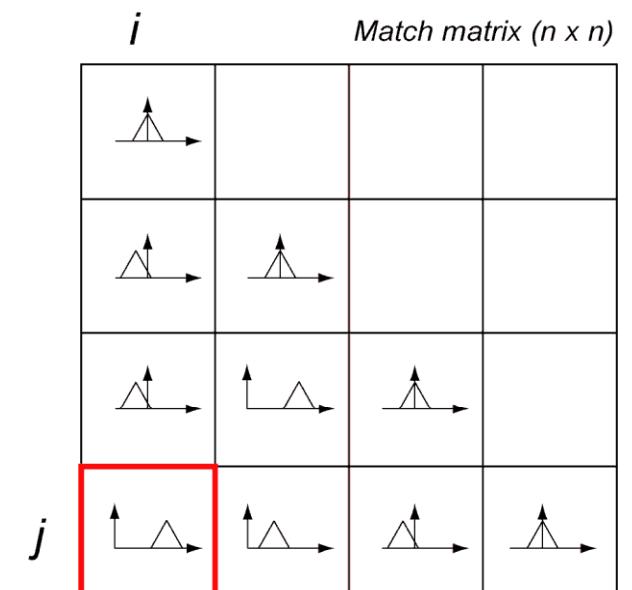
time $t' > t$



Rotation angle α : $\alpha = \tau(h_{ij}, \sigma_{ij})$

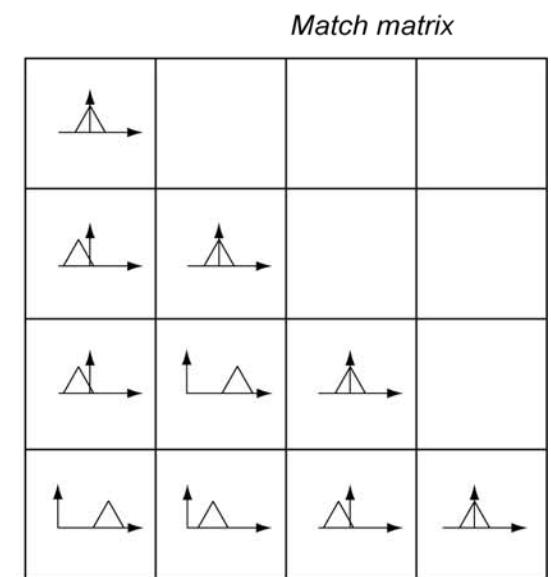
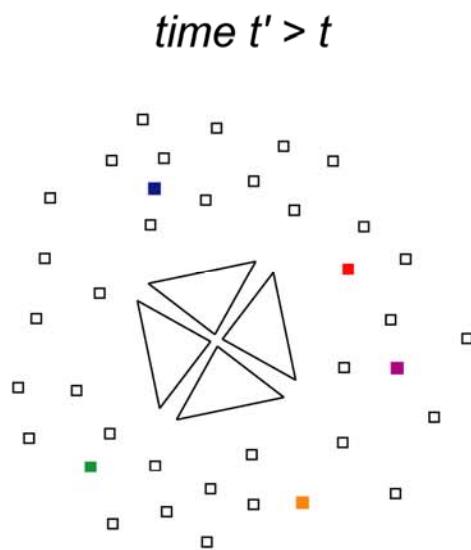
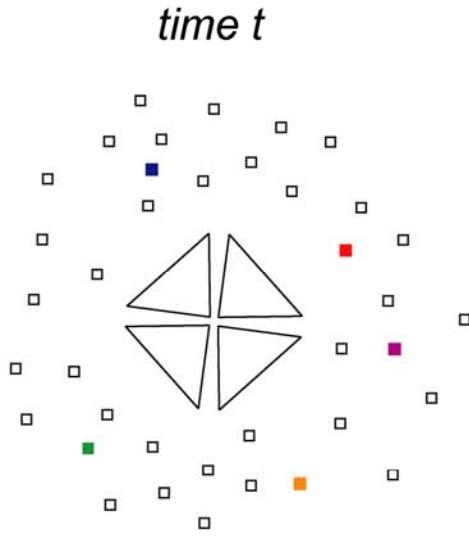
Estimate of α : $\bar{\alpha} = h_{ij}$

Error: $\varepsilon = \bar{\alpha} - \alpha = \tau(0, \sigma_{ij}) = \tau(0, \sigma)$



$\tau(h_{ij}, \sigma_{ij})$

Orientation from N correspondences



Estimate of α : $\bar{\alpha} = \frac{1}{N} \sum_{j=1}^N h_{ij}$ (N observations)

Error: $\varepsilon = \bar{\alpha} - \alpha = \frac{1}{N} \sum_{j=1}^N \tau(0, \sigma)$

$\xrightarrow[N \rightarrow \infty]{\text{CLT}}$ $\varepsilon \sim N(0, \sigma_N)$ $\sigma_N = \frac{\sigma}{\sqrt{N}}$

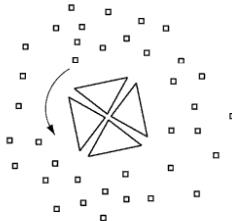
$f = 90^\circ$ $\sigma^2 = f^2/6$ $\sigma = 36.7^\circ$

$N = 200$ $\sigma_N = 2.6^\circ$

Rotation Guidance

Training

- Learn match matrix from training data
- Once and for all for a given camera configuration
- Arbitrary environment

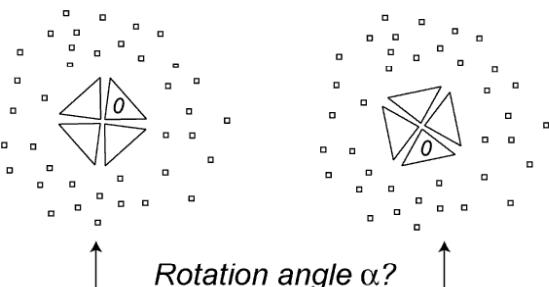


Match matrix ($n \times n$)

| | | | |
|--|--|--|--|
| | | | |
| | | | |
| | | | |
| | | | |

Navigation

- Use match matrix to estimate *relative* user orientation

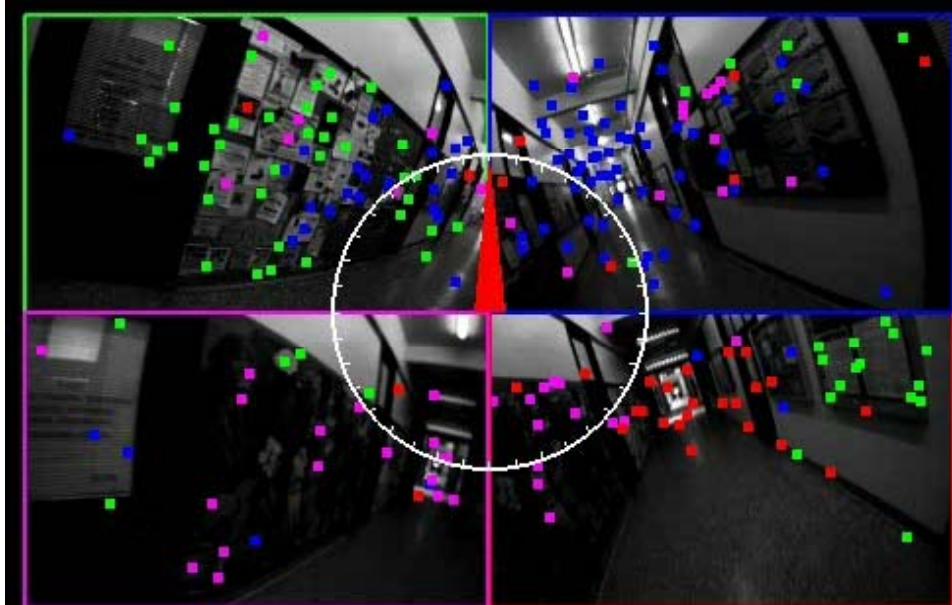


Match matrix ($n \times n$)

| | | | |
|--|--|--|--|
| | | | |
| | | | |
| | | | |
| | | | |

Rotation Guidance: Video

Live video stream



Localization



■ ■ ■ ■ Features

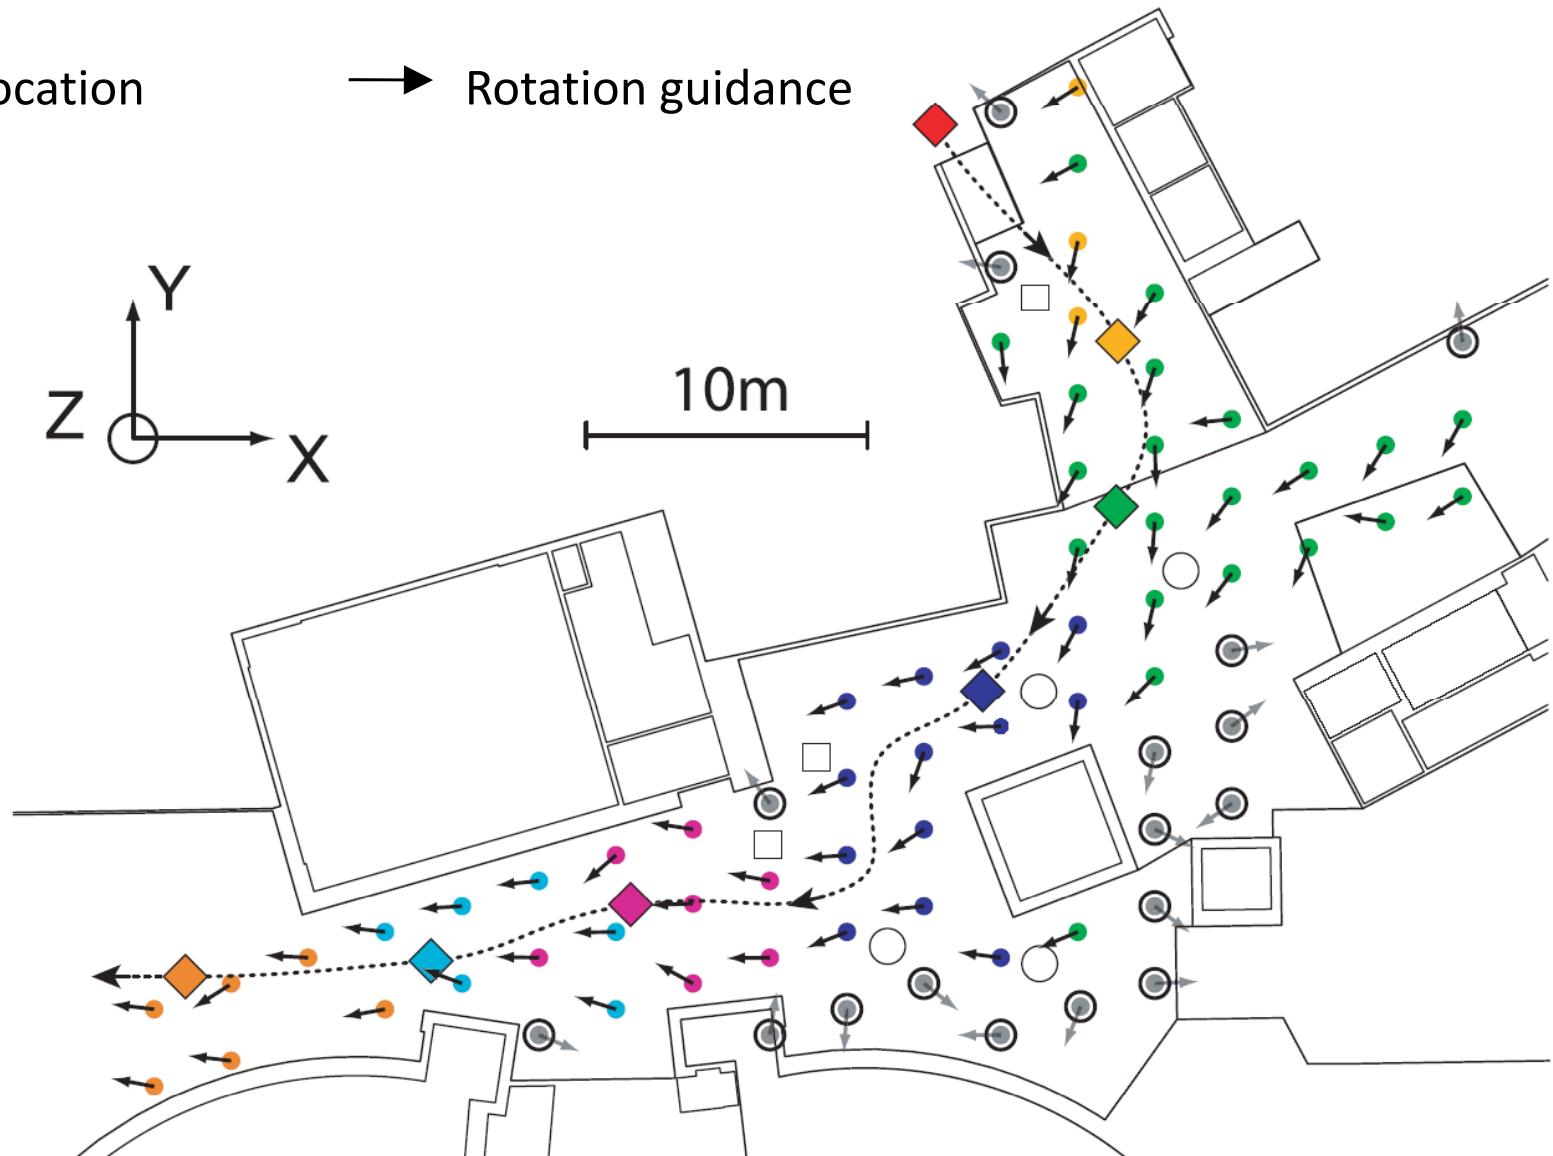
Rotation Guidance

◇ Place graph node

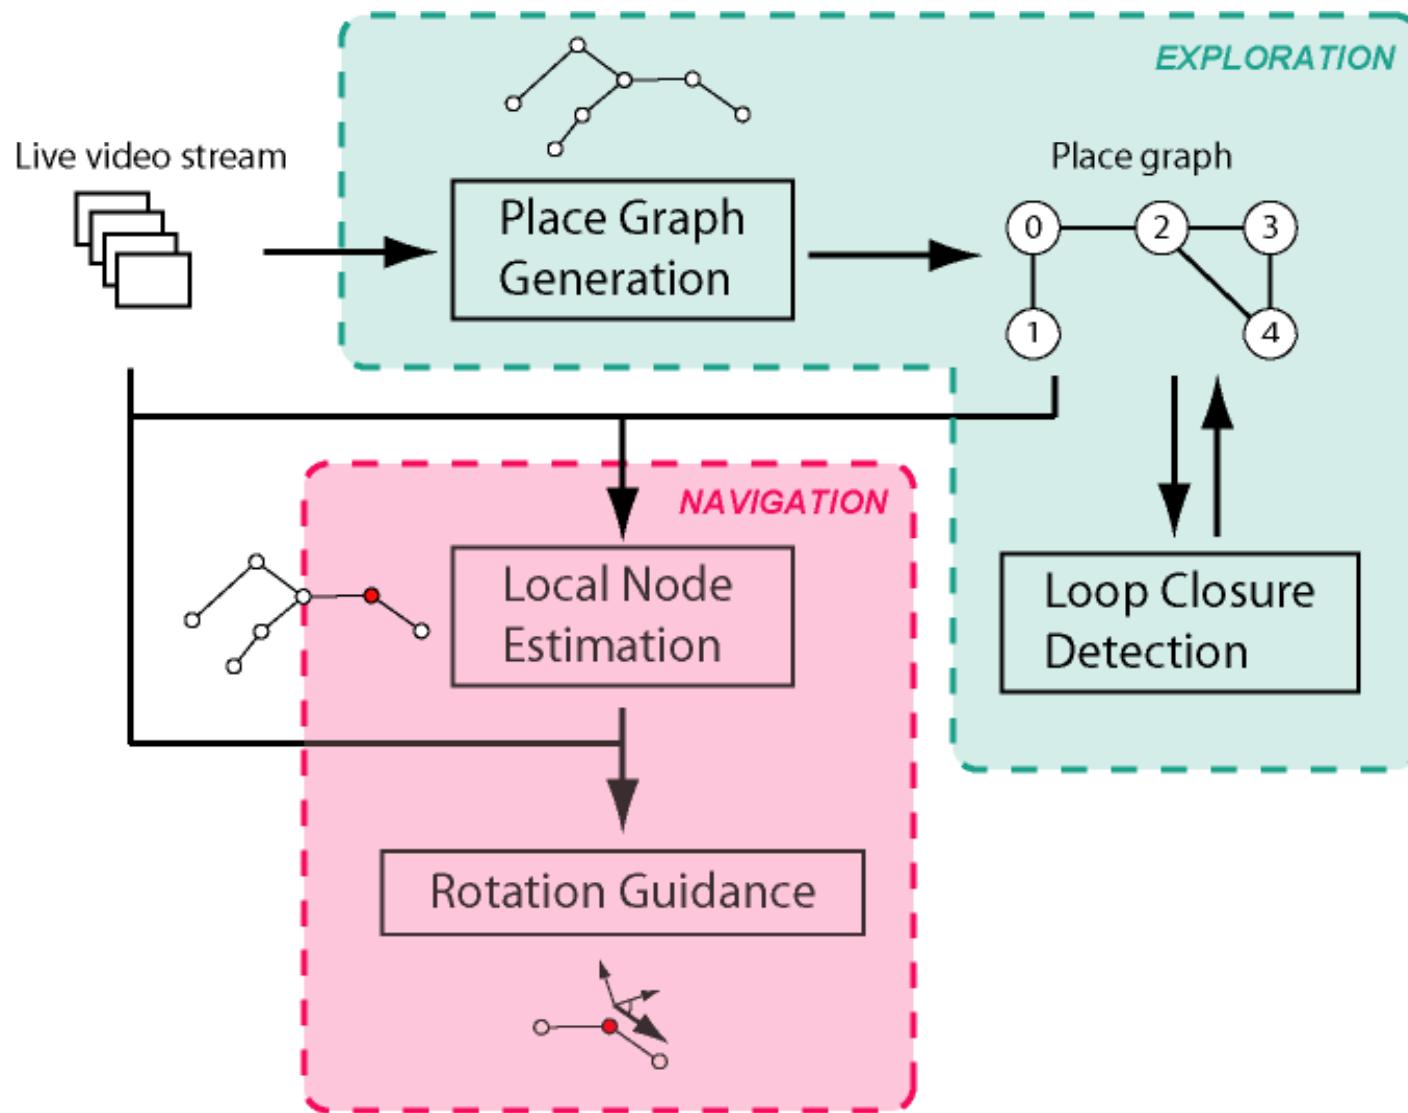
----- Exploration path (notional)

○ Test location

→ Rotation guidance

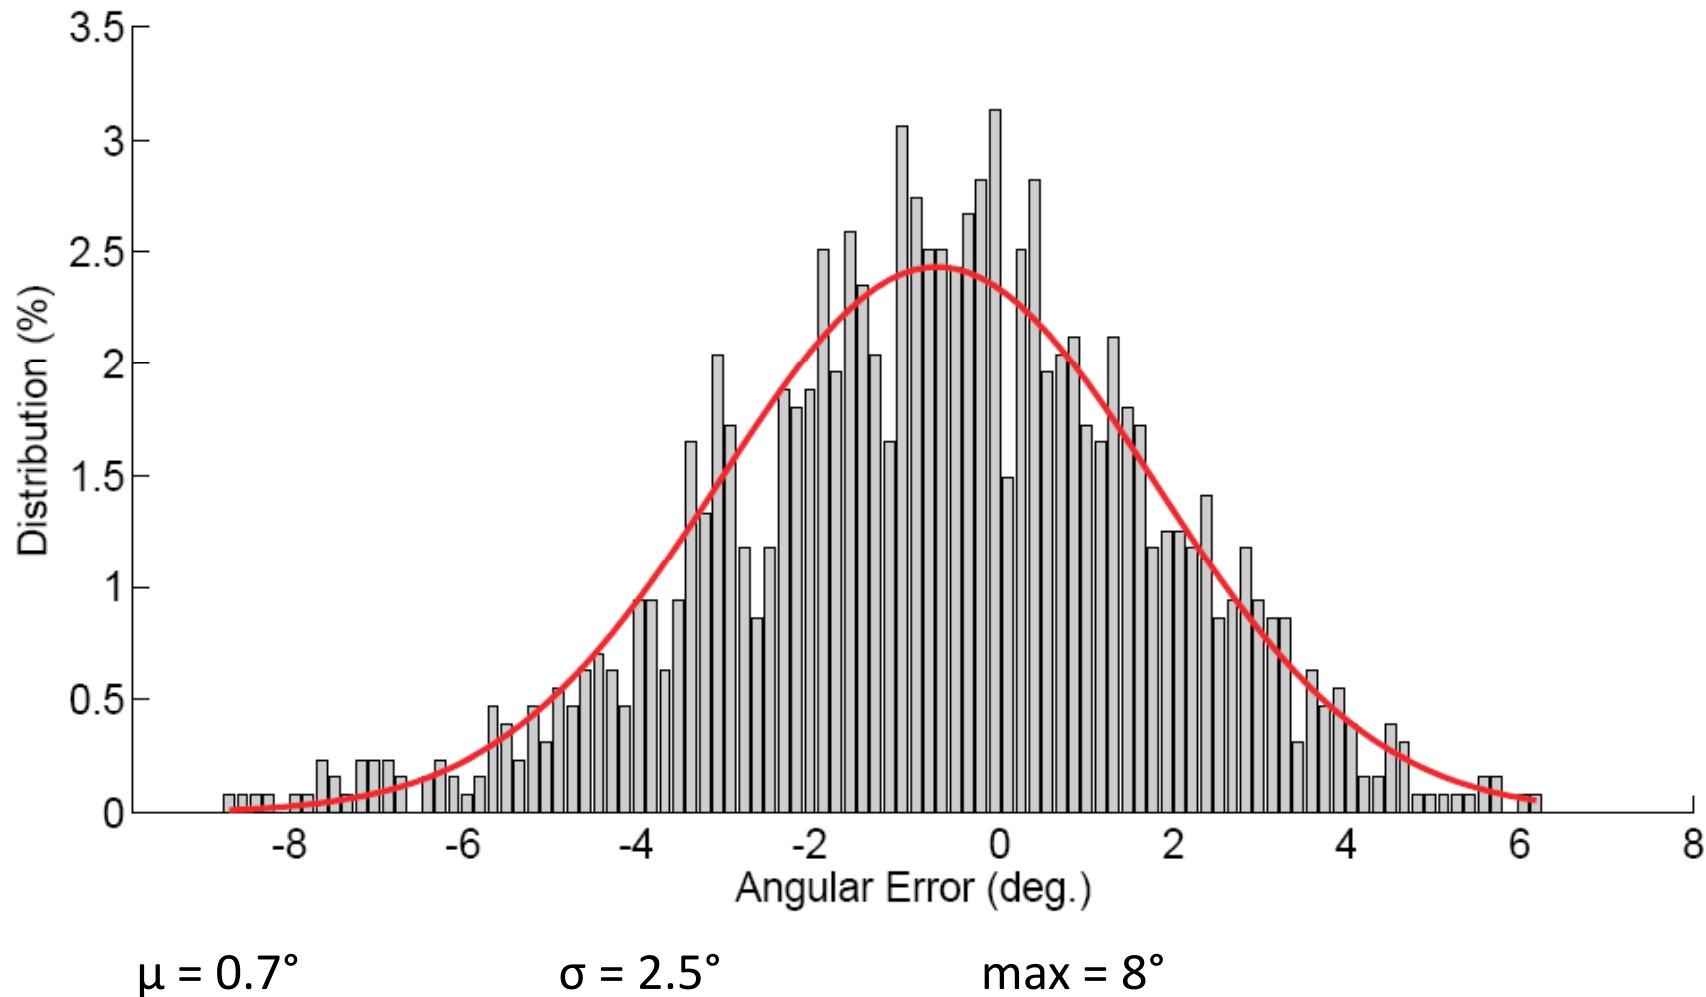


Method Overview



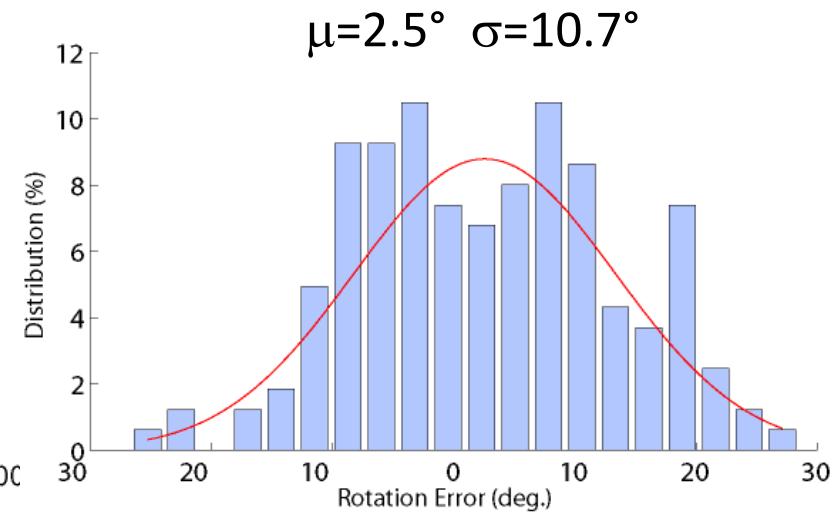
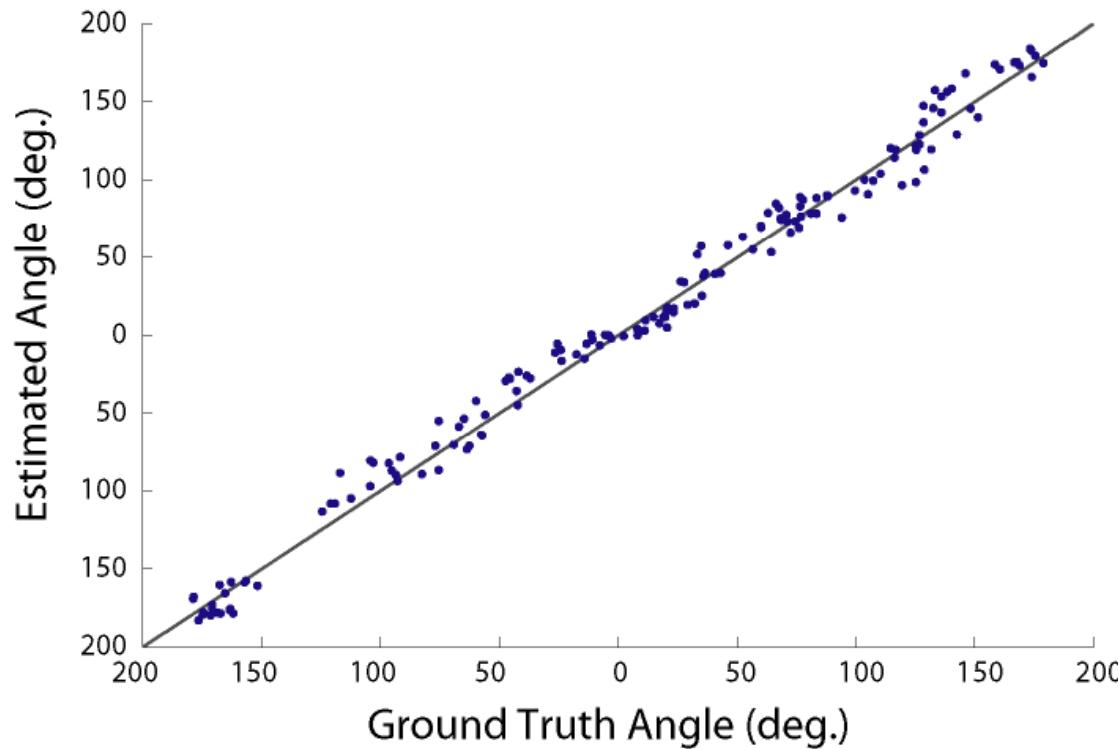
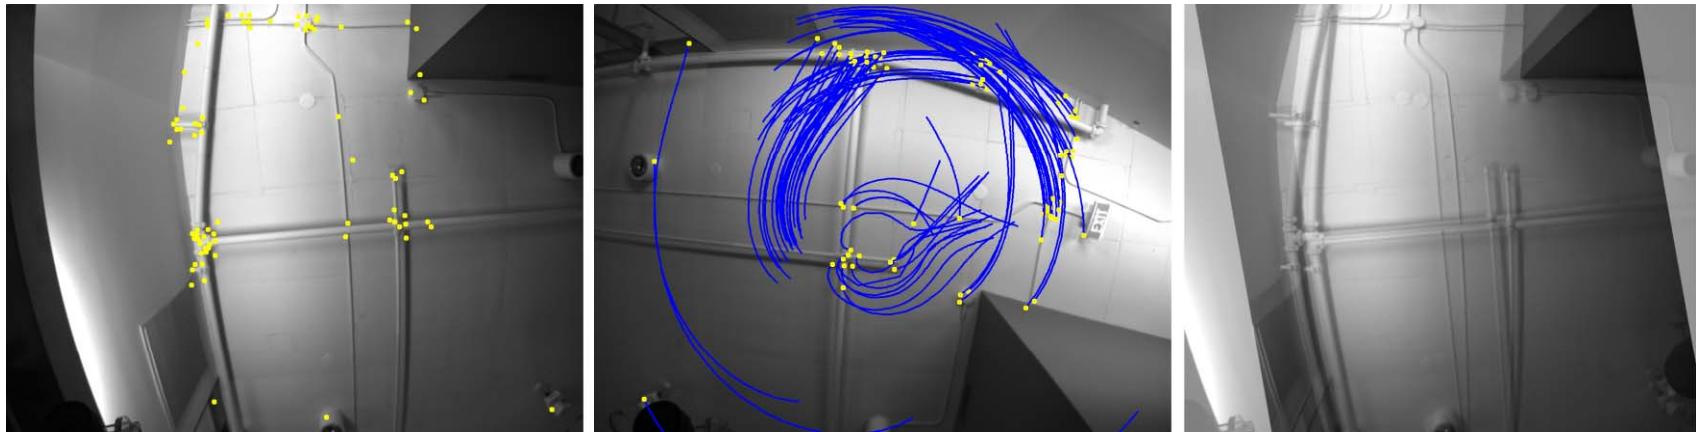
Method Evaluation

Rotation Guidance evaluation using IMU-derived ground truth (30 sec. sequence)



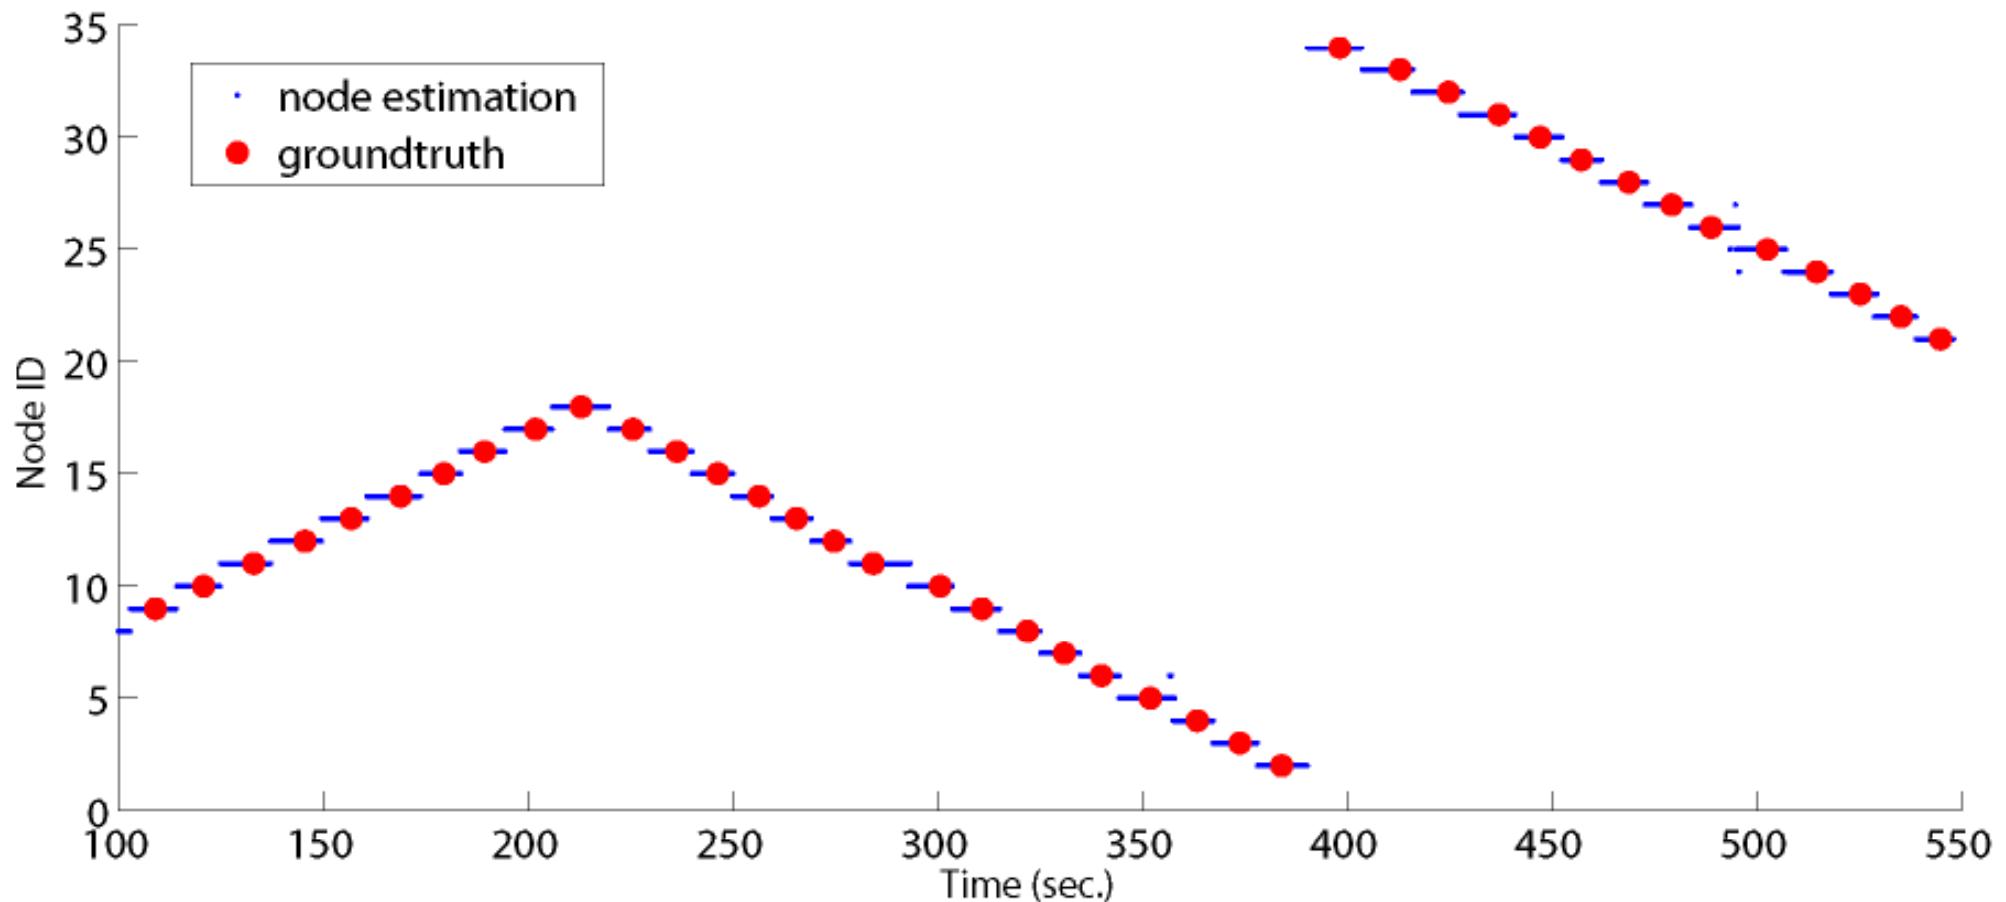
Method Evaluation

Large-scale rotation ground-truth using upward images

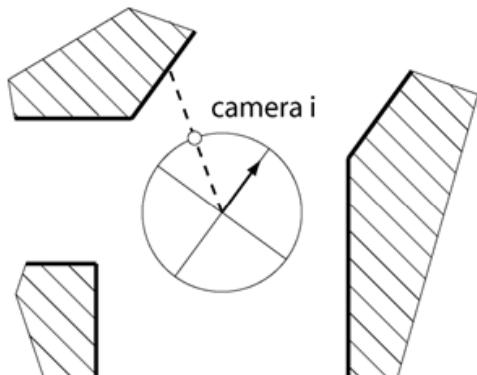


Method Evaluation

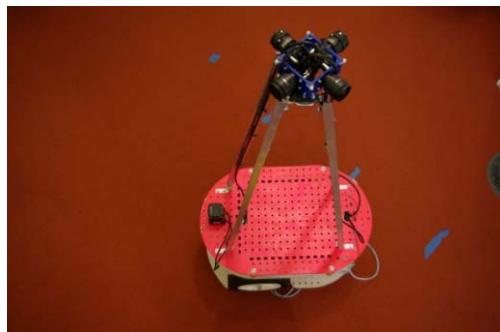
Large-scale localization ground-truth



Outline



- Sensor suite for vision-based navigation
- Body-relative navigation guidance using uncalibrated cameras
 - Topological mapping & loop closure
 - Localization & Rotation guidance
- Application to ground robot navigation

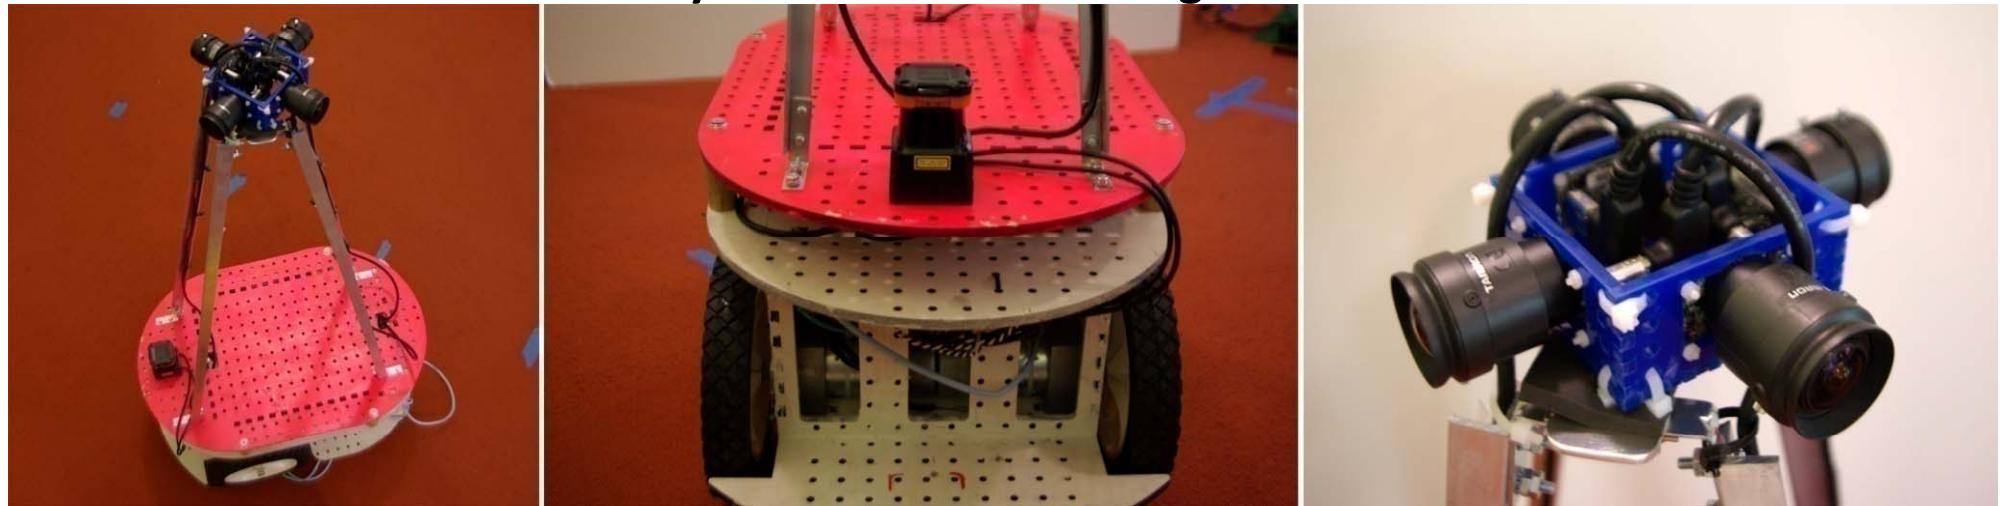


- User Study

Ground Robot Navigation

- Demonstrate approach on robot equipped with low-level obstacle avoidance capability
- Collect extended ground-truth data

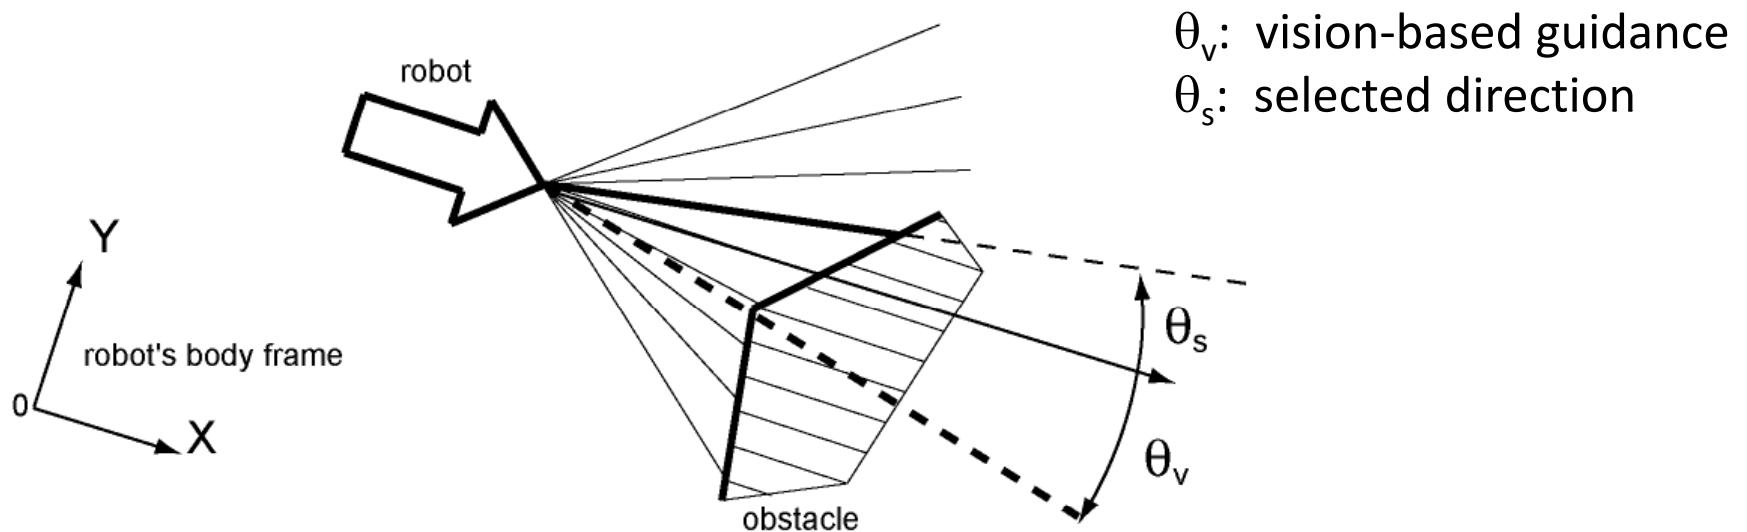
Two-wheel drive – IMU – Hokuyo LIDAR – 4-camera rig



Ground robot navigation using uncalibrated cameras, Koch, Walter, Huang, Teller, ICRA 2010

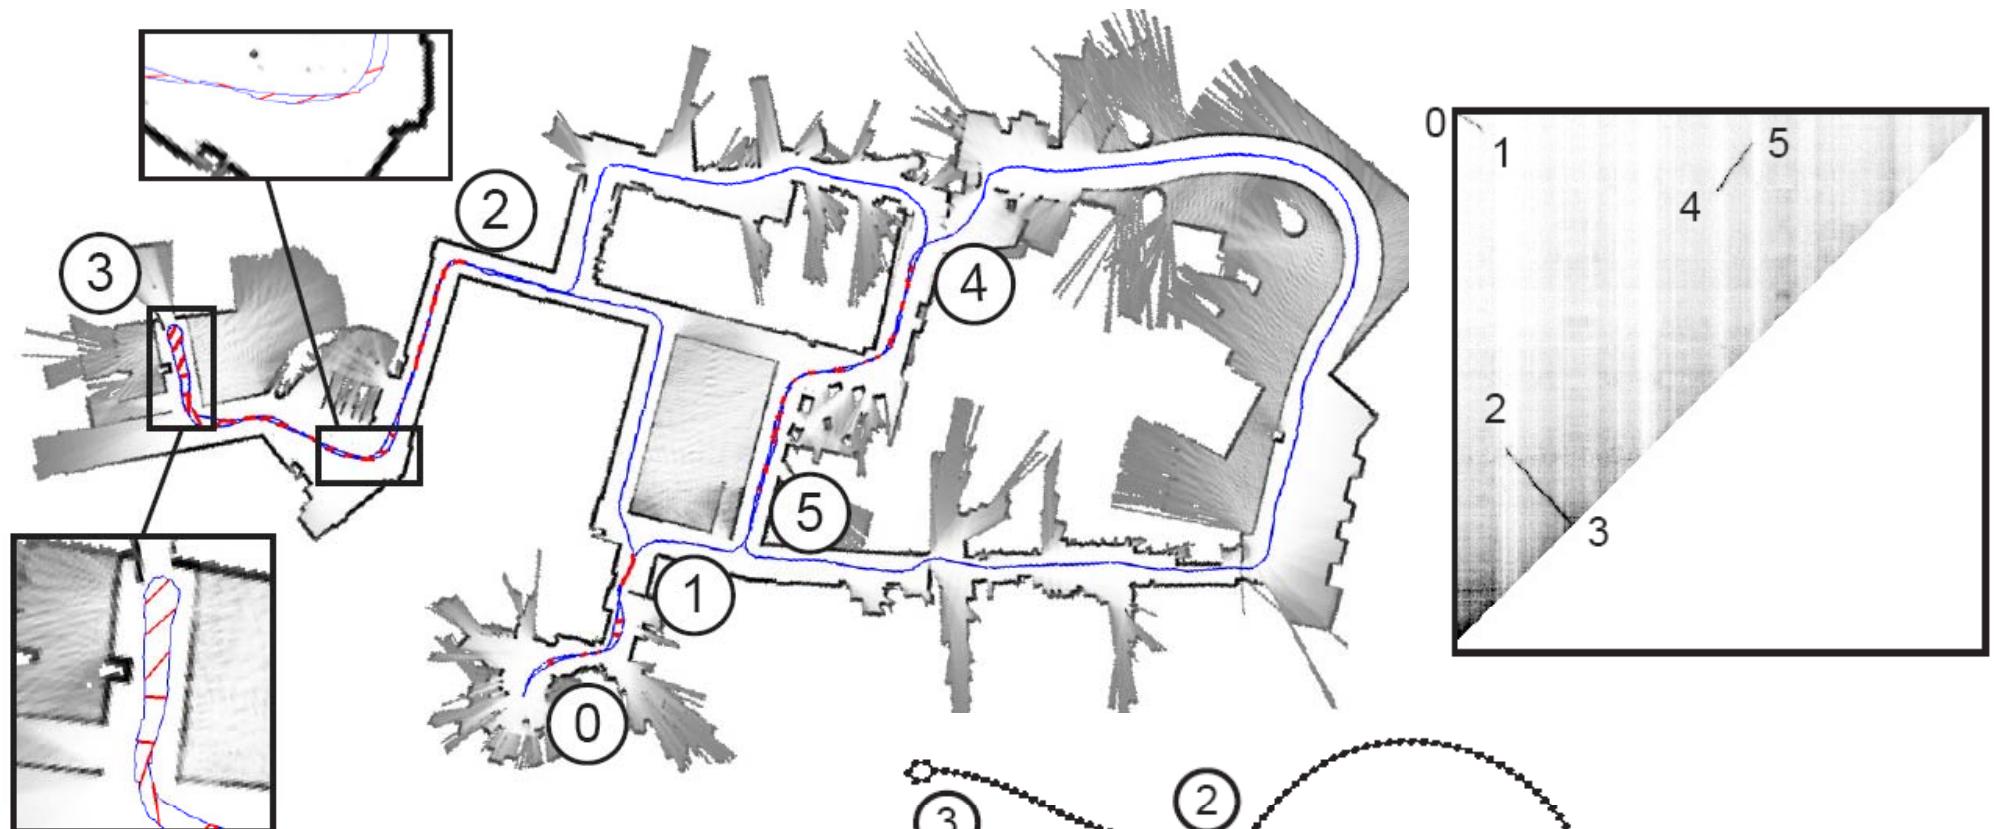
Ground Robot Navigation

1. Localization & high-level guidance using vision
2. “Greedy” local obstacle avoidance using laser



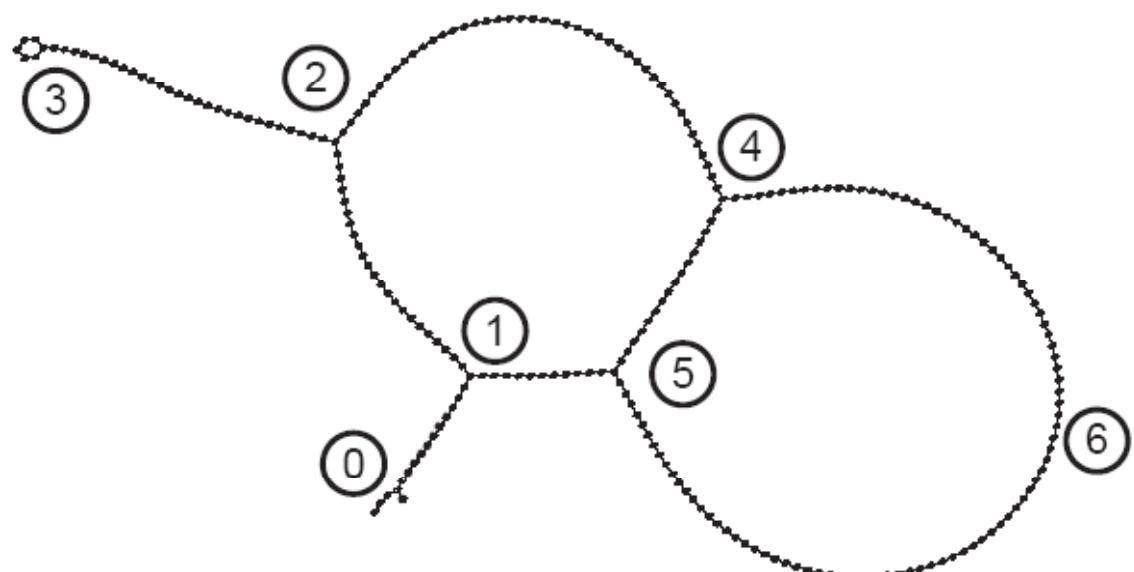
Ground robot navigation using uncalibrated cameras, Koch, Walter, Huang, Teller, ICRA 2010

Ground Robot Navigation



Loop closure – 24 minutes – 290 m

MIT Stata Center 3rd floor

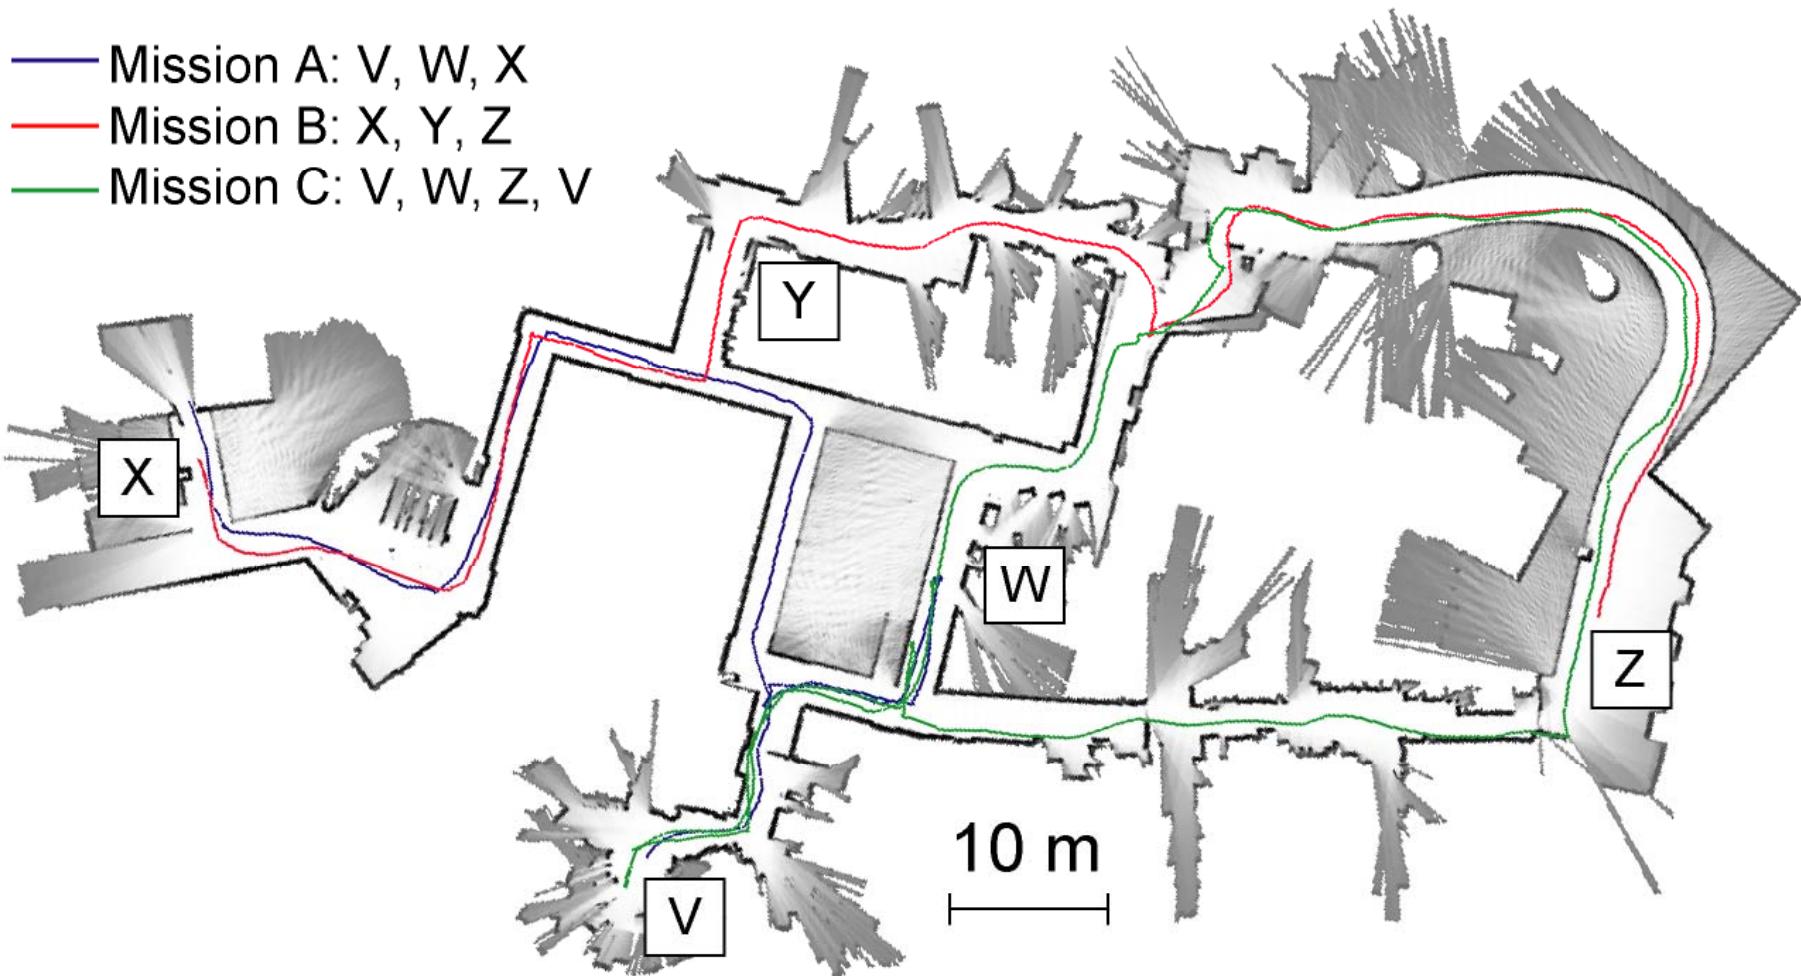


Ground Robot Navigation

3 missions 74 min – 380 m

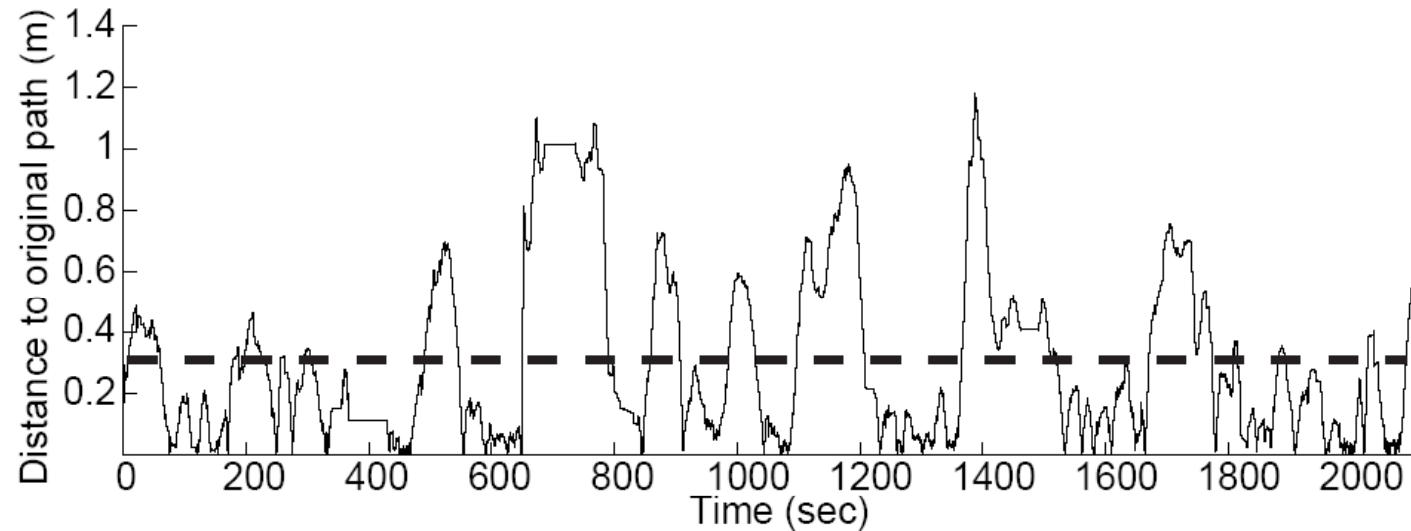
MIT Stata Center 3rd floor

- Mission A: V, W, X
- Mission B: X, Y, Z
- Mission C: V, W, Z, V

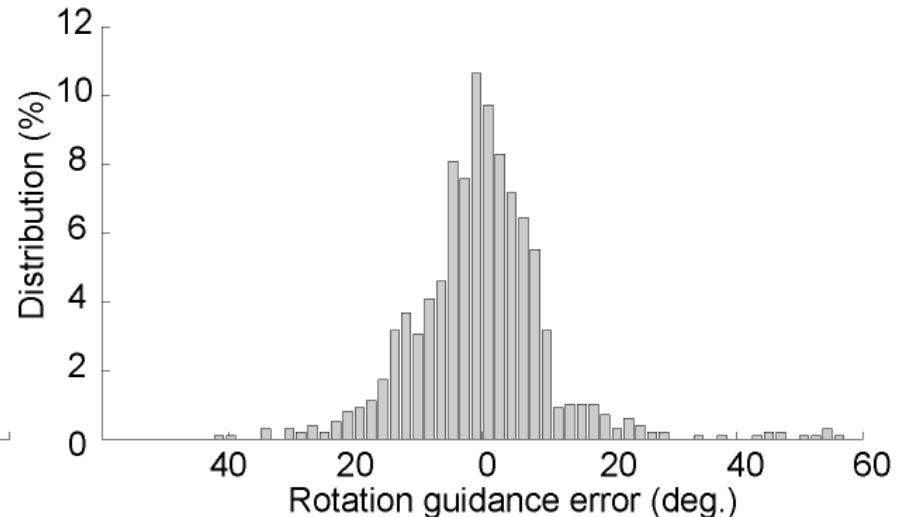
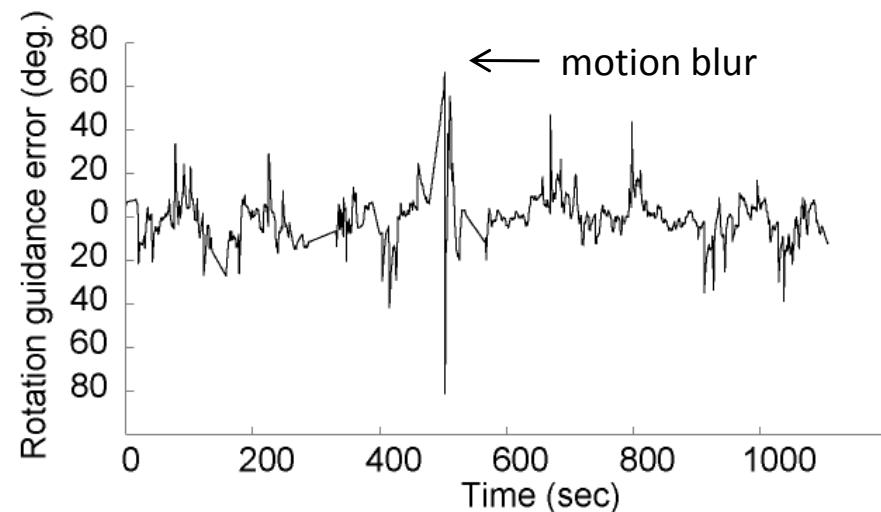


Ground Robot Navigation

Distance to original path (m)



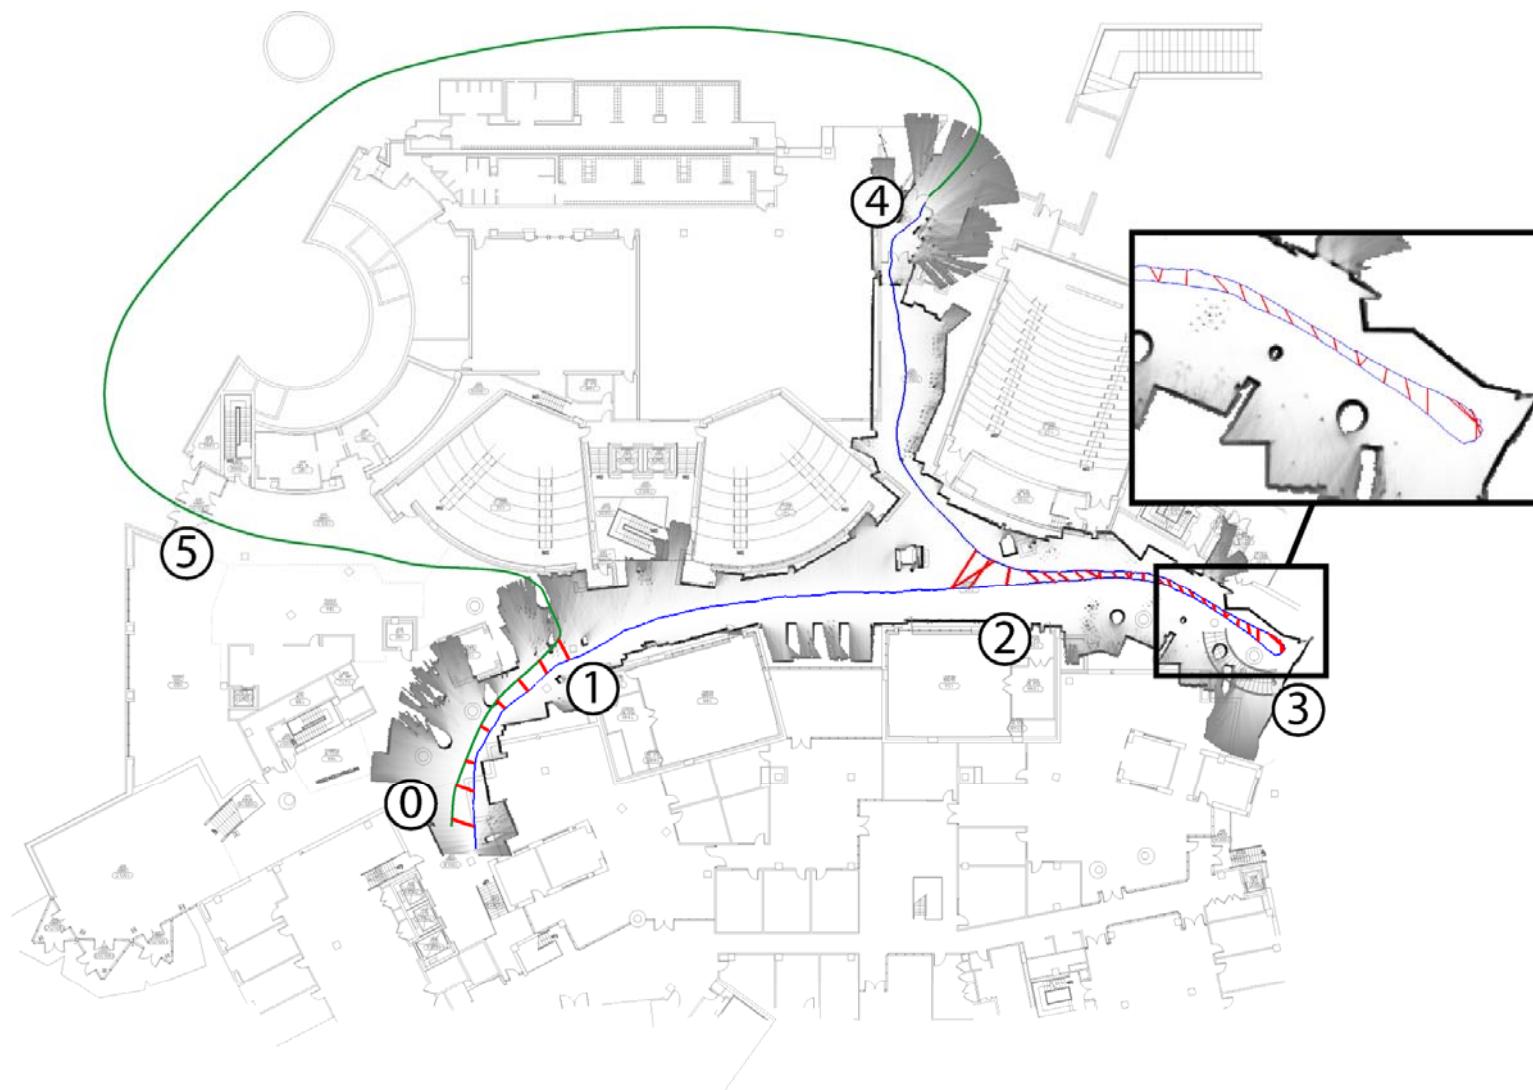
Rotation guidance error (deg.)



Ground Robot Navigation

Loop closure – 26 minutes – 400 m
MIT Stata Center 1st floor

- Outdoor path (notional, no mapping)
- Indoor path (ground-truth)
- Loop closure



Ground Robot Navigation Using Uncalibrated Cameras

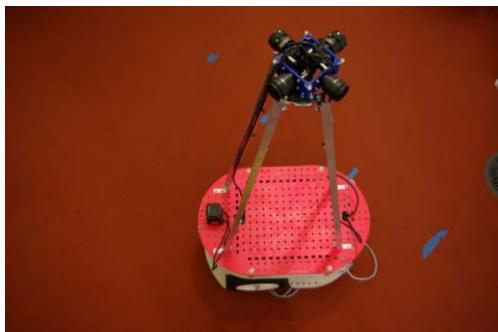
O. Koch, M. Walter, A. Huang, S. Teller

ICRA 2010

Outline



- Sensor suite for vision-based navigation
- Body-relative navigation guidance using uncalibrated cameras
 - Topological mapping & loop closure
 - Localization & Rotation guidance
- Application to ground robot navigation
- User Study



User Study

Experiment Setup

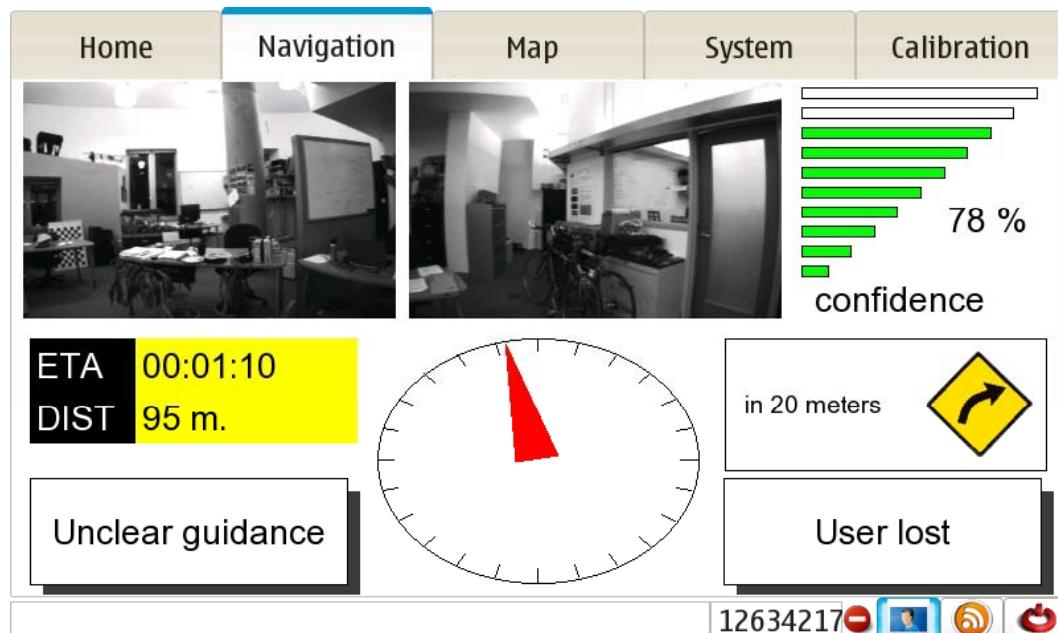
- Follow directions to a series of unknown target destinations in the building.

Effectiveness (qualitative metrics)

- User survey

Efficiency (quantitative metrics)

- # successful missions
- # unclear guidance events
- Speed ratio



Tablet PC User Interface



User Study

9 users, 2.5 hours, 6 km of exploration, 59 missions

Quantitative Analysis (*log data*)

| | Worst | Mean | Best |
|------------------|----------|--------|-------|
| Success rate | 50% | 70% | 100% |
| Speed ratio | 0.37 | 0.51 | 0.61 |
| Unclear guidance | 1.5 min. | 3 min. | Never |
| User lost | 3 min. | 5 min. | Never |

Qualitative Analysis (*survey*)

Positive Feedback

High-level guidance and non-metrical information are crucial.

Suggested improvements

More advanced turn-by-turn directions (Google Maps)

Compass/turn-by-turn inconsistency.

Drawing directions on images.

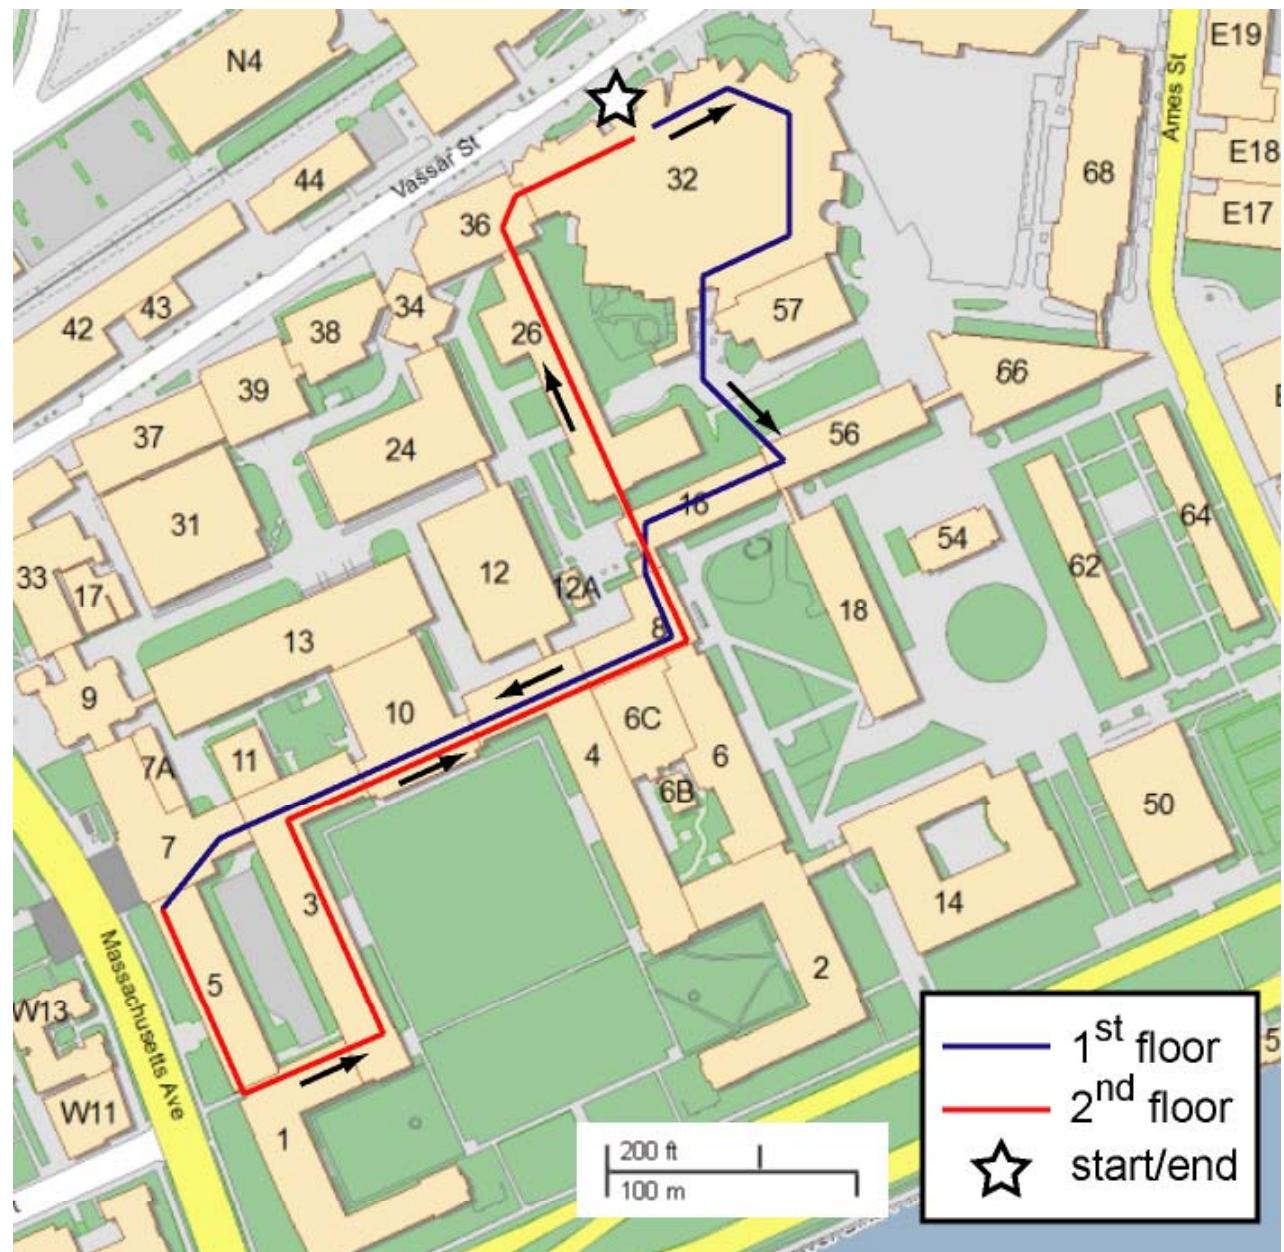
Simplify GUI.

User Study

Exploration Mission

23 minutes - 1. 1 km

Indoor/outdoor
Dynamic scenes



User Study

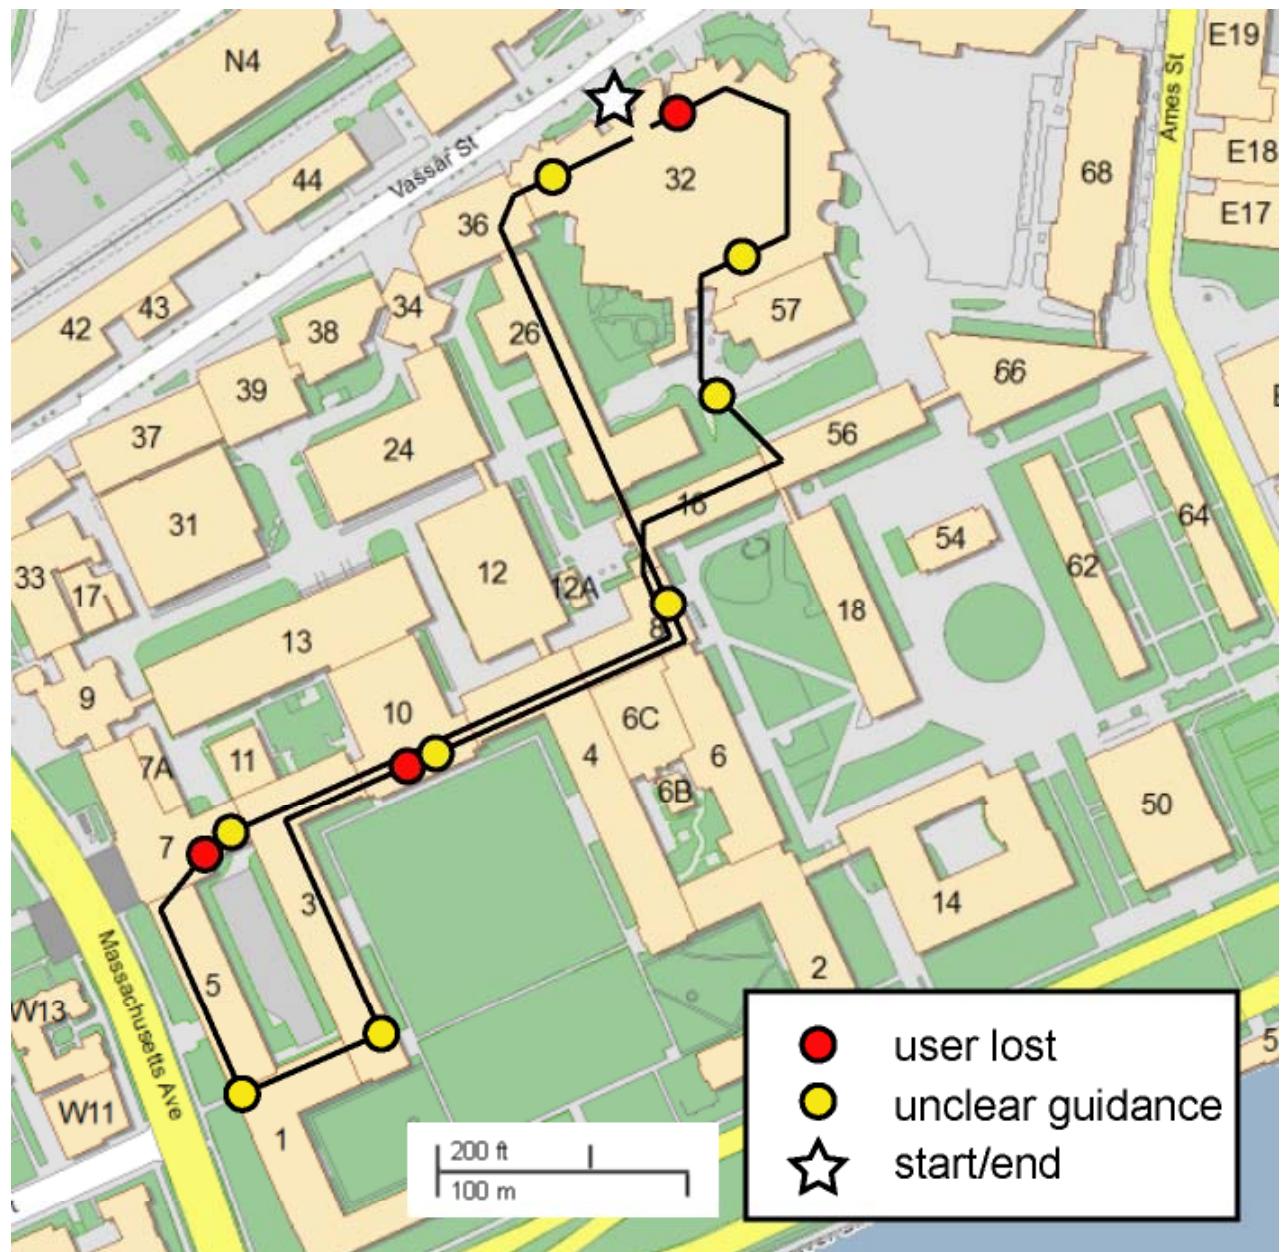
Navigation Mission

32 minutes

user lost: 3

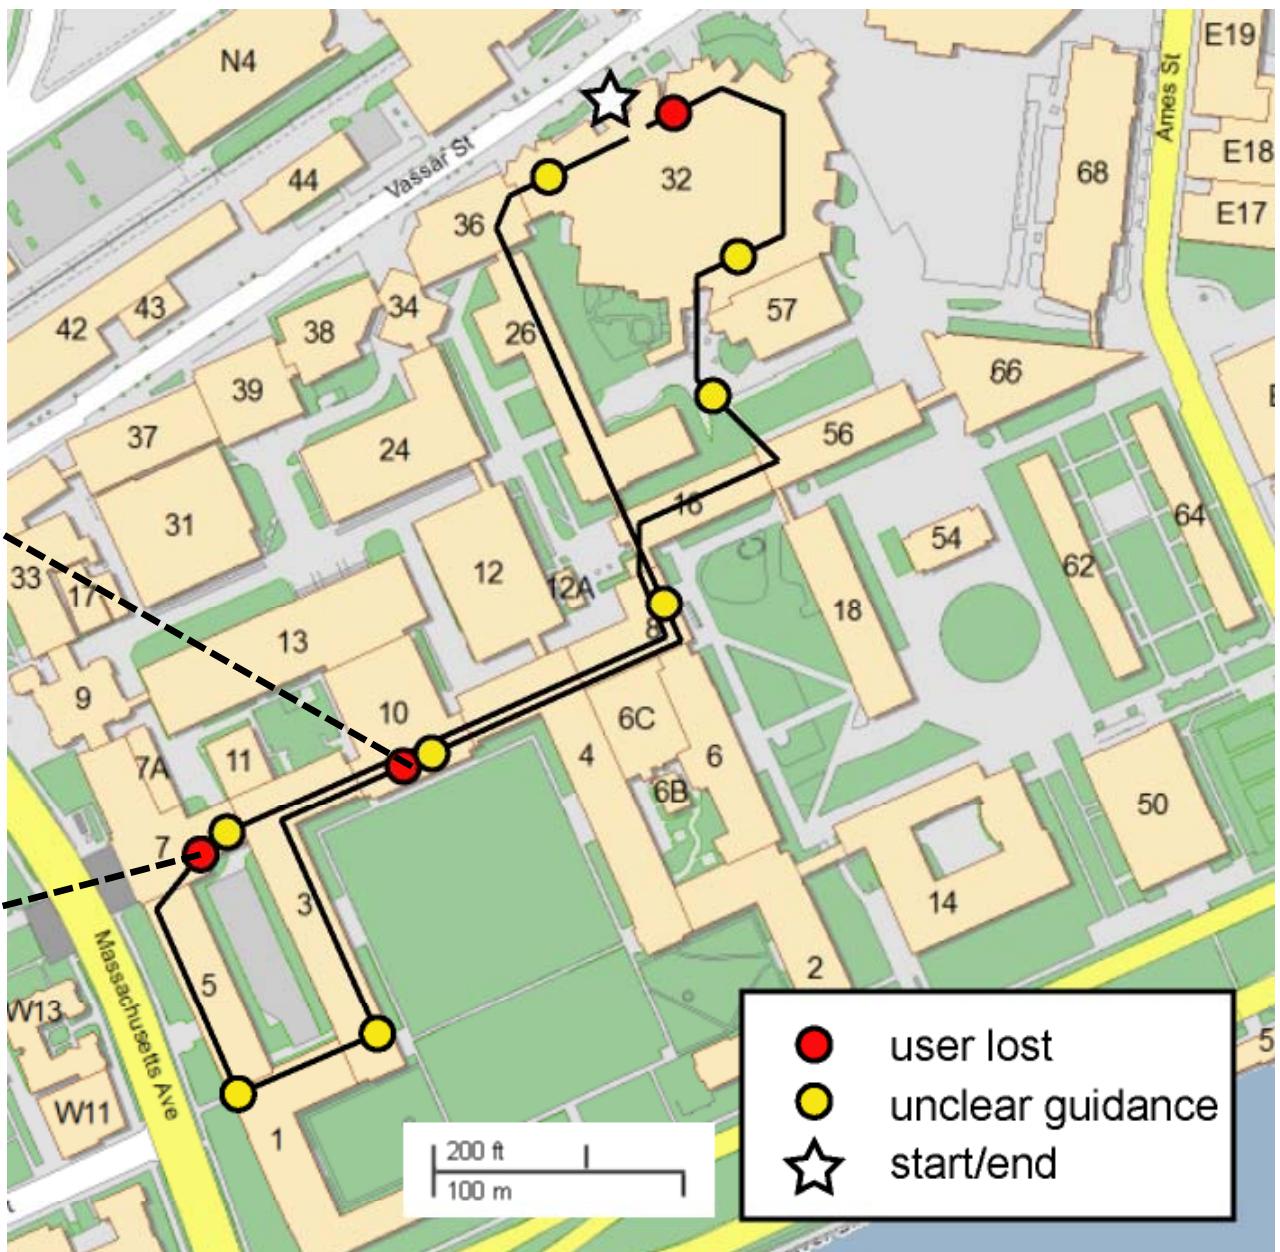
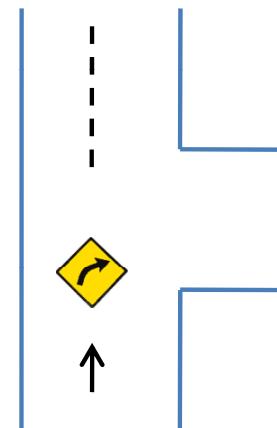
unclear guidance: 8

Speed ratio: 0.70



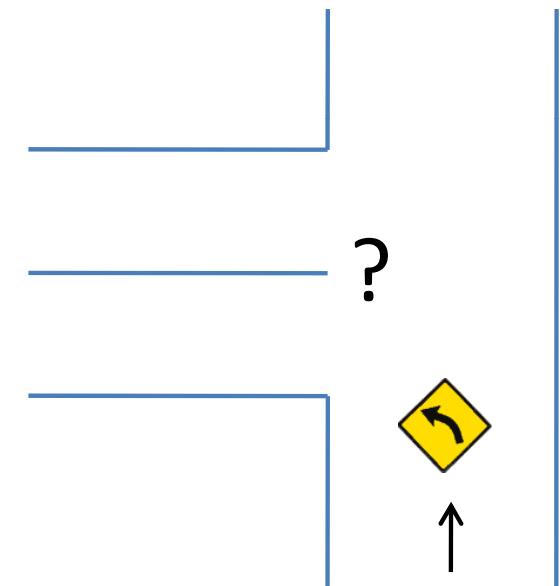
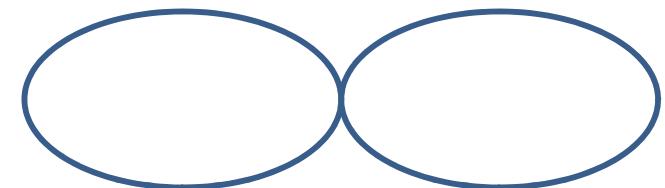
User Study

False positive

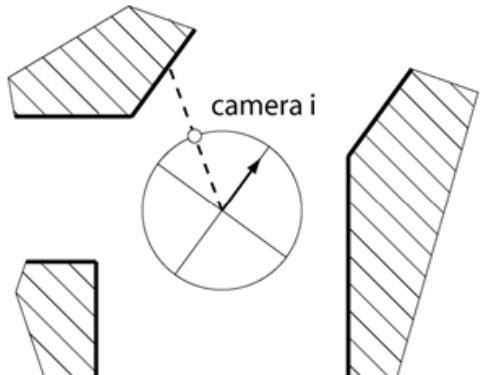


Failure Modes & Limitations

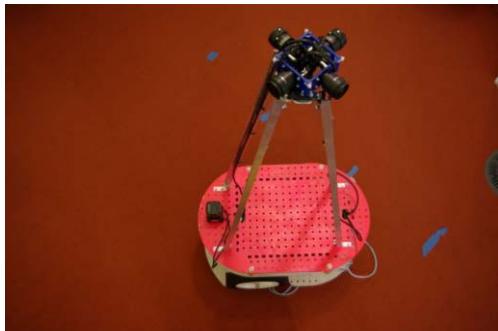
- Featureless environments
- Visually repetitive environments
- Non-isotropic feature distributions
- Ambiguous world configurations
- “One-shot” loop closure events



Contributions



- Sensor suite for vision-based navigation
- Body-relative navigation guidance using uncalibrated cameras
 - Topological mapping & loop closure
 - Localization & Rotation guidance
- Application to ground robot navigation
- User Study



Future Work

- Incorporate metrical constraints in the place graph
- Account for spatial relationship of features
- Bound expansion of visual vocabulary
- User interface
- Application to the visually impaired and in natural environments

Thank you

- Seth Teller
- Bill Freeman, Rob Miller
- The Draper Laboratory
- TIG, Bryt (Bradley), Ron (Wiken)

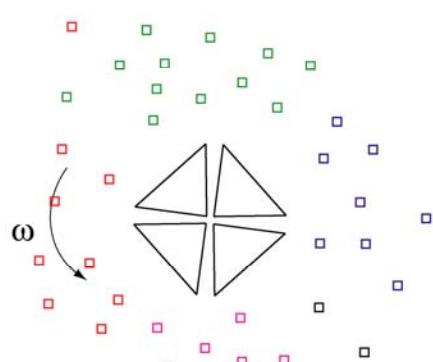


The 33x neighborhood

Questions?



Repeat for all features and all pairs of image frames (time sampling)



Match matrix ($n \times n$)

